Cost, Carbon, and the Speed of Replenishment in Distribution Network Design

Abstract

In this paper, we examine how a carbon price or cap changes two decisions of a firm that distributes through regional distribution centers: how many centers to operate and how fast to replenish them. Each decision trades one source of emissions against another. Slower line haul burns less fuel but requires more safety stock and more stock in transit; more distribution centers shorten the last mile but add buildings and give up the pooling that keeps inventory low. We model the two decisions jointly and show that the design problem has a unique optimum at every carbon price. For a given network, carbon pricing favors faster replenishment when the inventory that speed saves emits more per dollar than the speed itself. Goods in transit carry a financing cost but no carbon: for a given network, financing them raises the cost-minimizing speed without changing the carbon-minimizing speed. For a given speed, carbon pricing favors more distribution centers when a dollar of last-mile delivery emits more than a dollar of facilities and lost pooling. When speeds are interior, four carbon-intensity-per-dollar ratios and two cost shares of the cost-minimizing design give the direction of density in the joint design, and of speed when the two comparisons agree. The response of each decision to a carbon price depends on the difference between the carbon shares of the costs it trades off. A firm can compute these quantities from its operating rates, demand characteristics, and cost-minimizing configuration.

Keywords: Distribution network design; Carbon pricing; Replenishment speed; Safety stock; Last-mile delivery; Continuous approximation

1 Introduction

How many distribution centers to operate and how fast to replenish them are long-standing questions in distribution network design (Eppen 1979; Daskin et al. 2002; Shen et al. 2003), and firms continue to revisit both. Some have moved stock closer to demand. Amazon divided its US fulfillment network into eight regions and reported a 19% reduction in the miles traveled to deliver packages (Amazon.com, Inc. 2023), and Home Depot is building flatbed distribution centers in its forty largest markets (The Home Depot 2020). Others have slowed their line haul. AstraZeneca has moved most of its international freight from air to ocean and ships about 65% of its cargo by sea (Pharmaceutical Commerce 2023), and truckload shippers typically convert lanes to intermodal rail once the savings reach 10 to 15% (FreightWaves 2026). The reasons these firms give are mainly cost and delivery service; Amazon also describes shorter shipments as better for the environment, and AstraZeneca relates its shift to its sustainability efforts. What is new is that carbon is beginning to be priced and reported: the European Union has priced the carbon of maritime shipping since 2024 (European Parliament and Council of the European Union 2023) and will price road fuel from 2028 (European Environment Agency 2026). California is phasing in the reporting of supply-chain emissions from 2027 (State of California 2023), and ISO 14083 standardizes how the emissions of transport and warehouses are counted (International Organization for Standardization 2023). A firm that prices or caps its distribution carbon must therefore ask which way, and how far, a carbon price moves these two decisions.

In this paper, we examine how a carbon price changes these two decisions: the *density* of the network — how many distribution centers the firm operates — and the *speed* with which those centers are replenished. Each decision trades one source of carbon against another. Slowing replenishment moves freight by slower, lower-emission modes, as in slow steaming in shipping (Psaraftis and Kontovas 2013; Cariou 2011), but it lengthens the lead time that safety stock must cover, so the safety stock, and the carbon of holding it, rises with the square root of that lead time. Consolidating serves markets from fewer, larger distribution centers (McKinnon and Piecyk 2012) and pools demand across markets (Eppen 1979); it saves facility energy and inventory but lengthens the last mile. A firm thus provides *availability* either by holding safety stock or by replenishing faster, and market *proximity* either by operating more local facilities or by delivering farther from fewer. The cost-minimizing design uses one mix of these activities; with carbon priced, the mix can change.

The direction of the change depends on carbon intensity per dollar. For a given network, the carbon-minimizing chain replenishes *faster* when the inventory that faster replenishment saves emits more per dollar than the inbound transport that would replace it. For a firm that finances its goods in transit, faster replenishment saves capital as well as safety stock. Capital carries no carbon, so financing raises the cost-minimizing speed without changing the carbon-minimizing one. For a given speed, the carbon-minimizing chain operates a *more dispersed* network when last-mile travel emits more per dollar than the facilities that would shorten it and the pooling they give up.

We formalize this in a single-product location–inventory model with stochastic demand. The

firm chooses network density and replenishment speed to minimize cost subject to a cap on expected or realized carbon and a service requirement. Cost and carbon each comprise four components — facilities, holding, inbound transport, and last-mile transport — and cost also includes the financing of stock in transit. The model follows the location–inventory tradition (Daskin et al. 2002; Shen et al. 2003) in a continuous-approximation form (Daganzo 2005), with replenishment speed as a continuous decision. Because speed sets the lead time that safety stock must cover, it links transport and inventory in both cost and carbon, and the continuous form lets us state in closed form how pooling and the financing of goods in transit enter the comparison.

We contribute to the literature in three ways. First, we show how inventory pooling and the financing of goods in transit shape the cost–carbon comparison. The design problem has a unique optimum at every carbon price. For a given network, a carbon price favors faster replenishment only when the inventory that speed saves emits more per dollar than the speed itself, and financing goods in transit works against it. For a given speed, the comparison for density weighs the last mile against facilities and the pooling that an additional center gives up. When speeds are interior, four carbon intensities per dollar and two shares of the cost-minimizing design — the pooling weight, the share of a center's marginal cost that is lost pooling, and the safety-stock share, the share of the inventory that speed saves that is safety stock — give the direction of density, and of speed when the two comparisons agree. A price on transport carbon alone can raise total carbon, because it substitutes inventory and facilities for transport even where they emit more per dollar than that transport. Second, we estimate the ratios from public US data. Activities whose cost is mostly energy emit far more carbon per dollar than those whose cost is mostly labor or capital, so the direction depends on how the firm buys speed and how it delivers the last mile. At the central values, speed bought with fuel emits more per dollar than holding the stock of any product we examine, so a carbon price favors slower replenishment; diesel delivery calls for more distribution centers, and electric delivery for fewer. Third, we show what the services available on a lane and whole facilities mean for implementation. In a reference network that finances its stock in transit, slower line haul delivers nearly all of a 20% cut in emissions, changing both decisions costs less than slowing alone for deeper cuts, and the recommendation depends on the services available: with only truckload, intermodal, and carload rail, the least-cost design for a 30% cut moves to intermodal with slightly fewer centers, whereas with intermediate speeds it adds about one.

The paper is structured as follows. In Section 2, we position the paper relative to the location–inventory, carbon-network-design, and inventory–transport–emissions literatures. In Section 3, we develop the model, and in Section 4 we answer, in turn, the questions a firm asks of its network. Section 5 applies them to calibrated US data, Section 6 draws out the managerial implications, and Section 7 concludes.

2 Related Literature

Our work connects four streams of research: location–inventory design with risk pooling, carbon-aware network design, the joint control of inventory and transport mode under emissions regulation, and carbon as a priced cost in service-constrained inventory systems. Each develops part of the two substitutions we study — availability (inventory or replenishment speed) and proximity (facilities or last-mile travel). We draw on all four and add the joint design of network density and replenishment speed under a carbon cap and a service requirement, a classification of the direction in which carbon moves both decisions, and a calibration of the carbon intensities that determine it.

2.1 Location–inventory design and risk pooling

A first stream establishes how demand risk pooling shapes distribution network design: jointly choosing facility location and safety stock, it characterizes the cost-optimal number and placement of distribution centers. Eppen (1979) characterizes the effect of centralization on expected cost in the multi-location newsvendor, establishing the square-root relationship between the number of stocking locations and pooled safety stock. Daskin et al. (2002) and Shen et al. (2003) embed working and safety stock into the facility-location decision, so that risk pooling shapes the optimal number and placement of distribution centers, with extensions to stochastic settings by Snyder et al. (2007), to capacitated warehouses by Özsen et al. (2008), and to continuous-approximation allocation by Tsao et al. (2012). Daganzo (2005) provides the continuous-approximation foundation in which last-mile distance scales with the inverse square root of the number of facilities. Closest to our spatial–inventory coupling, Straubert (2024) ties the number of stocking locations to an exact reorder-point/order-quantity inventory model through a delivery-lead-time penalty, solved by continuous approximation. Recent work extends the location–inventory problem in directions close to ours. Hu et al. (2025) add the choice of storage temperature, trading cooling cost against the quality decay of fresh products in a network with risk pooling, and Jalal et al. (2026) integrate multi-period location, inventory, and transportation decisions under demand uncertainty for a pharmaceutical network. Puttemans et al. (2025) review this literature and call for more realistic modeling assumptions; our calibration in Section 5 anchors each carbon rate in public data. The centralize-versus-decentralize question these models pose can also depend on forces beyond demand pooling: Schmitt et al. (2015) show analytically that supply-disruption risk diversification can favor decentralization for a risk-averse firm, in contrast to the classical pooling result. These works develop network density and safety-stock placement under risk pooling with a cost or service objective. We retain the risk-pooling structure and the density decision and add a replenishment speed chosen together with density, and a carbon objective.

2.2 Carbon-aware network design

A second stream brings carbon into facility location and network design, characterizing how an emissions objective or cap can change the optimal network configuration. Cachon (2014) relates

retail store density to greenhouse gas cost, with consumer travel to stores as the dominant emission channel and inventory entering through retail floor space. Turken et al. (2017) examine facility location and capacity under a carbon tax and emission limits, weighing centralization against decentralization through production scale economies; Harris et al. (2011) compare cost-optimal and emissions-aware network designs numerically; Wang et al. (2020) develop a carbon-compliant location–inventory model under demand and carbon-price uncertainty; Li et al. (2017) jointly optimize distribution-network design and two-echelon inventory under a CO_2 emission tax; and Hacardiaux and Tancrez (2020) assess the transport-emission savings of a shared, consolidated location–inventory network under horizontal cooperation, with Hacardiaux et al. (2022) balancing partner preferences over logistics cost and carbon footprint in the same cooperative setting. Closest to ours in method, Zhang et al. (2026) combine continuous approximation with integer programming to show how a carbon price reshapes the choice among drone, truck–drone, and truck last-mile delivery. These studies make density or facility configuration the central decision and bring emissions into the objective, the constraints, or an ex post assessment; where transport is an operational choice, it enters through travel distances, discrete modes, or lead-time parameters. We treat network density and replenishment speed as joint decisions and characterize the proximity substitution between facilities and last-mile travel analytically.

2.3 Inventory, transport mode, and emissions

A third stream couples transport-mode and lead-time choices with inventory control under emissions regulation, characterizing how carbon reshapes the trade-off between freight and the inventory it serves. Hoen et al. (2014) and Chen and Wang (2016) study transport mode selection alongside inventory under stochastic demand and emissions regulation; Qi et al. (2022) select door-to-door transport services and routes for China–Europe freight against transport cost, in-transit inventory cost, and carbon; Bektaş and Laporte (2011) formalize the load–emission relationship in the pollution-routing problem, embedding fuel use and emissions as a function of vehicle load in the routing objective; Moradi et al. (2025) optimize driving speed and path selection under load- and grade-dependent fuel use at the routing level; Konur and Schaefer (2014) solve the economic order quantity model with less-than-truckload and truckload carriers under carbon cap, tax, and cap-and-trade regulation, and Schaefer and Konur (2015) extend this to a continuous-review, reorder-point inventory system under stochastic demand, characterizing the cost–emission Pareto frontier and analyzing lead-time and demand-variance sensitivity. Drent et al. (2023) analyze dynamic dual supply-mode selection under an emission target, and Hammami et al. (2015) bring carbon from manufacturing, ordering, and inventory into a multi-echelon production–inventory model with lead-time constraints. In shipping, Maloni et al. (2013) show by simulation that slower vessel speed lowers both total cost and carbon on an Asia–North America lane while raising the pipeline inventory that shippers carry. Our speed threshold prices the carbon as well as the cost of the inventory that slower replenishment requires, including the financing of stock in transit. In these papers, transport service enters through route and load physics, carrier class, discrete supply

modes, or lead-time parameters and constraints. We add the speed and density thresholds that result when replenishment speed and network density are chosen jointly under a carbon cap and a service requirement.

Closest in its combination of decisions, Comas Martí et al. (2015) design a supply chain network under uncertain demand in which each market is assigned a manufacturing facility and one of four discrete transport modes, with a continuous-review inventory policy at each market, the carrying cost of goods in transit, and caps or taxes on carbon. They solve integer programs numerically and show that functional and innovative products call for different networks and abatement, and that a carbon tax set at the marginal abatement cost achieves roughly the abatement of the corresponding cap. In their model each market holds its own stock, and lead time follows from the chosen mode and distance. We pool markets at distribution centers and treat replenishment speed as a continuous decision, which gives closed-form conditions for the direction of both decisions and shows how pooling and the financing of goods in transit enter them. Boone and Ganeshan (2025) choose a discrete freight mode against inventory under an emission cap at a single location. We replace the finite set of modes with a continuous replenishment speed, price carbon on the holding as well as the transport side, and embed the speed decision in a network whose density is chosen jointly.

2.4 Carbon policy, service, and safety-stock placement

A fourth stream prices carbon directly into inventory and service decisions, showing how a carbon price or cap shifts order quantities and safety-stock placement. Hua et al. (2011) and Bouchery et al. (2012) develop carbon-aware order-quantity models under cap-and-trade and multi-objective formulations, and Benjaafar et al. (2013) position cap, tax, and cap-and-trade as priced operational instruments in simple supply-chain models; Chen et al. (2013) derive the closed-form carbon-constrained economic order quantity, the single-location analog of the marginal carbon-intensity logic we extend to a network, and Arıkan and Jammernegg (2014) impose a product carbon-footprint constraint on a dual-sourcing newsvendor. Silbermayr et al. (2017) bring environmental constraints into inventory pooling with dependent demand. Within multi-echelon service models, Ni and Shu (2015) study the trade-off between service time and carbon in safety-stock placement under the guaranteed-service framework, and Wang et al. (2025) incorporate cap-and-trade into guaranteed-service placement. Probabilistic carbon limits have been modeled through chance constraints in lot sizing (Koca 2021) and in supply-chain network design (Shaw et al. 2016), typically inside large mixed-integer programs; our deterministic equivalent scales the transport component of carbon by a single factor. A recent review (Sebatjane 2025) surveys carbon-aware inventory models and the regulatory instruments they treat. The service-placement models price carbon against service and safety stock; we bring network density and replenishment speed into the same comparison.

In summary, prior work covers carbon-aware density, mode and lead-time choice, service-constrained placement, and chance constraints on carbon. We couple replenishment speed with network density under carbon and service constraints, show that the joint design is unique at every

carbon price, and classify the carbon-minimizing direction of both decisions. We then estimate from public data the carbon intensities that decide that direction.

3 The Model

3.1 Operational environment

Region and markets. A firm serves N markets distributed over a service region of area A. Per-period demand at market i is a random variable with mean μ and standard deviation σ, identically distributed across markets with common pairwise correlation $\rho \in [0, 1)$. Demand is stationary and the setting is single-product. Markets and distribution centers are treated symmetrically, the standard simplification that admits closed-form continuous-approximation results (Daganzo 2005).

Distribution centers and geography. The firm opens n identical distribution centers (DCs), $1 \le n \le N$, each serving a sub-region of area A/n. By the continuous-approximation spatial law, the mean last-mile (DC-to-market) distance falls as the square root of the served area,

$$d(n) \;=\; \delta\, n^{-1/2}, \qquad \delta \propto \sqrt{A}, \tag{1}$$

so a denser network shortens delivery but only at a decreasing rate (Daganzo 2005).

Replenishment and the speed decision. Each DC is replenished directly from the source at a chosen speed $v \in [v_{\min}, v_{\max}]$. Individual DCs lie at different distances from the source; what enters cost and carbon is total inbound freight: with identical DCs each carrying throughput $N\mu/n$, this is $\sum_j (N\mu/n)\, D_j = N\mu\, \bar{D}$, the total throughput times the *mean* source-to-DC distance $\bar{D}$. Because the DCs fill a region of fixed size served from a fixed source, $\bar{D}$ is set by the region's geometry (order $\sqrt{A}$) and is approximately invariant to n; we set $D \equiv \bar{D}$ and treat it as constant.[1] The replenishment lead time is then

$$L(v) \;=\; D/v. \tag{2}$$

Lead time is the only channel through which speed affects inventory: safety stock covers demand over the lead time, and the goods in transit are proportional to it. Shortening the lead time is not free. We use $t(v)$ and $\varepsilon(v)$ to denote the cost and carbon of inbound freight per unit-mile at speed v, and we assume only that both increase with v and are convex; for the closed-form results, we use

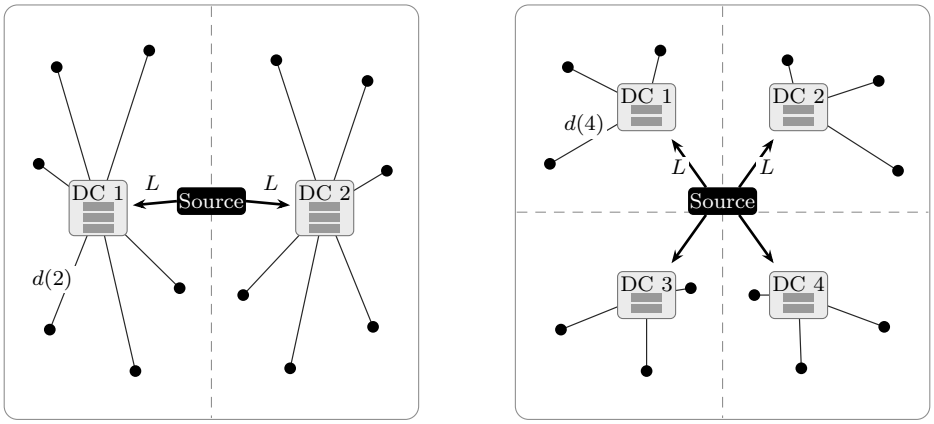
[1] Two modeling reductions are embedded here. First, the symmetric-DC reduction assigns every DC the representative lead time $L = \bar{D}/v$ in place of its own $L_j = D_j/v$. Safety stock enters through $\sqrt{L}$, which is concave in distance, so by Jensen's inequality $\sqrt{\bar{D}} \ge n^{-1} \sum_j \sqrt{D_j}$ (Jensen 1906): using the mean distance slightly *overstates* aggregate safety stock, by a gap that is second-order in the dispersion of the D_j and reduces as the distances equalize. Second, replenishment is point-to-point; were the DCs served on a consolidated route, the inbound distance would grow as $\sqrt{n}$ in the manner of the store-delivery tour of Cachon (2014), adding a term that reinforces centralization in Proposition 4.3 and leaving the fixed-density speed comparison of Proposition 4.2 unchanged; with inbound distance varying in n, however, the speed response would depend on density through $D(n)$, so the complementarity of Lemma 4.1(iii) need not carry over unchanged.

the quadratic instantiation $\varepsilon(v) = \varepsilon_0 v^2$ and $t(v) = t_0 + t_1 v^2$, the canonical form motivated by the cubic fuel-speed law of slow steaming (Psaraftis and Kontovas 2013; Cariou 2011). The model does not represent how a shorter lead time is achieved; it represents only the cost and carbon of achieving it. On a given origin–destination lane, a faster mode such as truckload instead of intermodal rail, a faster vessel, or a team of two drivers who keep the truck moving all deliver a shorter lead time at a higher freight cost, and they differ in how much of that cost is fuel.

Fuel and other costs of replenishment. We divide the transport cost rate into fuel and other costs, $t(v) = t^{\mathrm{f}}(v) + t^{\mathrm{o}}(v)$. With a fixed fuel price and emission factor, and with fuel accounting for the emissions that vary with speed, $\varepsilon'(v) = \chi_{\mathrm{f}}\, t^{\mathrm{f}\prime}(v)$, where χ_{f} is the carbon per dollar of fuel, so that

$$\frac{\varepsilon'(v)}{t'(v)} = \chi_{\mathrm{f}}\, \xi(v), \qquad \xi(v) = \frac{t^{\mathrm{f}\prime}(v)}{t'(v)}. \tag{3}$$

When both cost components rise with speed, $\xi(v)$ is the fuel share of the marginal cost of speed. Under the quadratic instantiation, $t_1 = b_{\mathrm{f}} + b_{\mathrm{o}}$ and $\varepsilon_0 = \chi_{\mathrm{f}} b_{\mathrm{f}}$, where $b_{\mathrm{f}} v^2$ and $b_{\mathrm{o}} v^2$ are the speed-dependent fuel and other costs, so that ε_0 / t_1 is constant. Moving from intermodal rail to truckload shortens lead time by burning more fuel per unit-mile, so most of its added cost is fuel. A team of two drivers shortens lead time at nearly the same fuel, and so nearly the same carbon, per mile, so most of its added cost is the second driver's wages. Where t is a freight rate, the relation assumes that fuel costs pass through to the rate.



$n = 2$: each DC serves 6 markets
more stock per DC, longer delivery legs

$n = 4$: each DC serves 3 markets
less stock per DC, shorter delivery legs

Faster replenishment (shorter L), at a given density:
system-wide safety stock shrinks, so holding cost and carbon fall, while freight cost and carbon rise.

More DCs (higher n), at a given speed:
delivery legs $d(n)$ shorten with territory area, so system-wide last-mile cost and carbon fall, while facility cost and carbon rise and weaker pooling raises system-wide holding cost and carbon.

Figure 1: Model representation for 12 markets (N) and 2 or 4 DCs ($n = 2$ or 4).

Inventory and service. Each DC operates a base-stock policy and must meet a cycle service level α — the in-stock probability at the end of a period, the Type-1 (cycle) service measure. With normally distributed demand the safety factor is $z = \Phi^{-1}(\alpha)$, with $\alpha > \frac{1}{2}$ so that $z > 0$, taken here as a service requirement (a parameter). Demand is independent across periods, and each DC reviews its inventory continuously, so safety stock protects against demand over the replenishment lead time; under periodic review the protection interval would also include the review period. A DC that pools $g = N/n$ markets faces aggregate per-period demand with standard deviation $\sigma\sqrt{g\,[1 + (g-1)\rho]}$, the statistical economy of scale in inventory (Eppen 1979). Covering lead-time demand, each DC holds safety stock $z\,\sigma\sqrt{L}\,\sqrt{g[1 + (g-1)\rho]}$. Summing over the n DCs and defining the *pooling expression*

$$\Psi(n) \;=\; N\rho + n(1-\rho), \tag{4}$$

the system-wide safety stock is

$$\mathrm{SS}(n, v) \;=\; z\,\sigma\,\sqrt{D/v}\,\sqrt{N\,\Psi(n)}. \tag{5}$$

At $\rho = 0$, (5) reduces to $z\sigma\sqrt{L}\sqrt{nN}$, the square-root law: decentralizing from 1 to N DCs multiplies safety stock by $\sqrt{N}$. As $\rho \to 1$, $\Psi(n) \to N$ and safety stock becomes invariant to n — the pooling benefit vanishes.

Decisions. The firm chooses the network density $n \in [1, N]$ (relaxed to a continuous variable) and the replenishment speed $v \in [v_{\min}, v_{\max}]$. Service is a requirement, so α (hence z) is a parameter in our model. Table 1 collects the notation.

3.2 Cost and carbon

Both the cost and carbon functions are system-wide flows over a single representative period, consistent with the stationary single-product environment above: F, φ are charged per DC per period, h, η per unit held per period, and the transport terms accrue against the per-period throughput. The representative distances in Figure 1 therefore enter global cost and carbon through aggregate freight tasks, not through a per-territory objective. Throughput itself is conserved at $N\mu$ regardless of the configuration: the firm moves the mean demand of all N markets through the network, whatever density n and speed v it chooses. With per-period throughput $N\mu$,

$$K(n, v) = \underbrace{F\,n}_{\text{fixed}} + \underbrace{h\,\mathrm{SS}(n, v)}_{\text{holding}} + \underbrace{h_p\,P(v)}_{\text{in transit}} + \underbrace{t(v)\,N\mu D}_{\text{inbound}} + \underbrace{t_L\,N\mu\,d(n)}_{\text{last-mile}}, \tag{6}$$

$$C(n, v) = \underbrace{\varphi\,n}_{\text{fixed}} + \underbrace{\eta\,\mathrm{SS}(n, v)}_{\text{holding}} + \underbrace{\varepsilon(v)\,N\mu D}_{\text{inbound}} + \underbrace{\varepsilon_L\,N\mu\,d(n)}_{\text{last-mile}}, \tag{7}$$

where F, φ are fixed cost/carbon per DC, h, η are holding cost/carbon per unit per period, h_p is the financing cost of a unit in transit per period, $P(v) = N\mu D/v$ is the stock in transit, and t_L, ε_L

Table 1: Notation. The decision variables are the network density n and the replenishment speed v; all other quantities are exogenous parameters or derived.

Symbol	**Meaning**
Decision variables	
n	number of DCs (network density), $1 \le n \le N$
v	replenishment speed, $v \in [v_{\min}, v_{\max}]$
Demand (exogenous)	
N	number of markets
$\mu,\ \sigma$	mean, standard deviation of demand per market per period
ρ	cross-market demand correlation, $\rho \in [0, 1)$
Geography, cost, carbon (exogenous)	
$A,\ D,\ \delta$	region area; inbound distance; last-mile distance constant
$F,\ \varphi$	fixed cost, carbon per DC per period
$h,\ \eta$	holding cost, carbon per unit per period
h_p	financing cost of stock in transit per unit per period
$t(v),\ \varepsilon(v)$	inbound transport cost, carbon per unit-mile (convex in v)
$t_L,\ \varepsilon_L$	last-mile transport cost, carbon per unit-mile
α	cycle service level (in-stock prob.); $z = \Phi^{-1}(\alpha)$
$\tau,\ B$	carbon price; carbon cap
Derived	
$g = N/n$	markets served per DC
$L = D/v$	replenishment lead time
$d(n) = \delta\, n^{-1/2}$	mean last-mile distance
$\Psi(n) = N\rho + n(1-\rho)$	pooling expression
$\mathrm{SS}(n, v)$	system-wide safety stock, eq. (5)
$P(v) = N\mu D/v$	stock in transit
$a(n, v)$	safety-stock share of the inventory saving from speed, eq. (8)

are last-mile cost/carbon rates (Shen et al. 2003; Hua et al. 2011; Benjaafar et al. 2013). Each term multiplies a rate by the activity it prices; apart from the financing of stock in transit, which carries no carbon, the two functions differ only in whether the rate is measured in dollars or kilograms.

Facilities. The facility terms Fn and φn are the periodic cost and carbon of operating the open DCs. They rise linearly in n and do not depend on v. Facility carbon φ is distinct from holding carbon η: φ captures baseline building energy and embodied carbon, whereas η prices the incremental carbon of storing a unit (for example refrigeration or handling).

Holding. The holding terms $h\mathrm{SS}(n, v)$ and $\eta\mathrm{SS}(n, v)$ value the system-wide safety stock of (5). They rise in n through the loss of pooling in $\Psi(n)$ and fall in v through the lead-time factor $\sqrt{D/v}$. This is the channel through which replenishment speed couples transport and inventory carbon; for cost, stock in transit is a second channel.

Stock in transit. Each unit shipped is in transit for $L = D/v$ periods, so an average of $P(v) = N\mu D/v$ units are in transit, and this stock falls as replenishment speeds up. A firm that owns the

goods while they travel carries their financing cost, h_p per unit per period — the capital charge on the value of the goods; $h_p = 0$ when title passes on receipt at the distribution center. Stock in transit occupies no warehouse, and the carbon of moving it is counted in the inbound transport term, so it adds no carbon of its own; its financing cost does not depend on n. Faster replenishment therefore saves two kinds of inventory: safety stock, which carries holding carbon, and stock in transit, which carries none. We use

$$a(n, v) = \frac{h\,\mathrm{SS}(n, v)}{h\,\mathrm{SS}(n, v) + 2h_p\,P(v)} \in (0, 1], \tag{8}$$

to denote the share of the inventory saving from speed that is safety stock; the factor two arises because safety stock falls with the square root of the lead time and stock in transit with the lead time itself. Without financing, $a = 1$; with financing, the share rises with speed.

Inbound transport. The inbound terms multiply the representative freight task $N\mu D$ by the speed-dependent rates $t(v)$ and $\varepsilon(v)$. Since D is treated as invariant to n, inbound transport does not vary with network density; it rises in v through the increasing convex rates $t(v)$ and $\varepsilon(v)$, opposing the holding term that falls in v.

Last-mile transport. The last-mile terms multiply throughput $N\mu$ by the mean DC-to-market distance $d(n) = \delta n^{-1/2}$ and the rates t_L, ε_L. A denser network shortens this leg at the continuous-approximation rate, so the last-mile term is decreasing and convex in n and independent of v under constant per-mile last-mile rates.

Decomposition and scope. The carbon function splits into a transport component $\varepsilon(v)N\mu D + \varepsilon_L N\mu d(n)$ and a facility-and-holding component $\varphi n + \eta \mathrm{SS}(n, v)$; the cost function has the analogous partition, with the financing term $h_p P(v)$ in addition. Slowing replenishment reduces inbound transport carbon but raises safety stock, while dispersion reduces last-mile distance but raises facility count and pooling loss. We hold the replenishment quantity of the network fixed, so that each center's order falls as centers are added and system-wide cycle stock is invariant to both decisions, and we exclude production and procurement because throughput is fixed at $N\mu$.

3.3 The program

The firm minimizes cost subject to a carbon cap B and the service requirement,

$$\min_{1 \le n \le N,\ v_{\min} \le v \le v_{\max}} K(n, v) \quad \text{s.t.} \quad C(n, v) \le B. \tag{9}$$

Here K and C are expected flows over the model's accounting period: under the stationary demand assumption they are the long-run average cost and carbon of a configuration, and they can be scaled to an annual expected-carbon cap by multiplying by the number of periods. Because demand is random, realized carbon over a specified accounting horizon fluctuates about its expectation; we

develop the analysis for the expected-value cap and return below to a chance-constrained cap that bounds realized carbon with specified confidence (Corollary 4.6).

Equivalently, pricing carbon at $\tau \geq 0$ gives the priced objective $J = K + \tau C$. At an interior optimum, the cap multiplier equals the price ($\lambda = \tau$), as in carbon-constrained inventory models (Chen et al. 2013); sweeping τ traces the cost–carbon frontier, along which the optimal configuration moves continuously (Lemma 4.1). Because cost and carbon are convex in the logarithms of n and v (Lemma 4.1), the capped problem (9) is a convex program in those variables: every cap above the carbon minimum is met by a finite carbon price, and the capped problem and the problem under that price share the same unique optimum. The carbon minimum itself is the limit as the price grows without bound. Jumps arise only from the integer number of distribution centers. Define the blended rates

$$F_\tau = F + \tau\varphi, \qquad H_\tau = h + \tau\eta, \qquad T_\tau(v) = t(v) + \tau\varepsilon(v), \qquad \ell_\tau = t_L + \tau\varepsilon_L.$$

Then

$$J(n, v; \tau) \;=\; F_\tau n + H_\tau \,\mathrm{SS}(n, v) + h_p\, P(v) + T_\tau(v) N\mu D + \ell_\tau N\mu d(n). \tag{10}$$

The financing cost of stock in transit carries no carbon and is not blended. Setting $\tau = 0$ recovers cost; the limit $\tau \to \infty$ recovers carbon up to scale.

4 Analysis

We now characterize how a carbon price changes the two decisions. The answers depend on four carbon intensities per dollar, which we use throughout: holding, $r_h = \eta/h$; inbound transport at the margin, ε'/t' of (3), which equals $r_t = \varepsilon_0/t_1$ under the quadratic instantiation; last-mile delivery, $r_L = \varepsilon_L/t_L$; and facilities, $r_F = \varphi/F$. Each measures the carbon that a dollar spent on the activity emits, so the direction in which carbon moves a decision depends on each activity's carbon *relative to its cost*. We take up, in turn, the questions a firm asks of its network: whether its design responds to a carbon price in a predictable way (Section 4.1), when replenishment should be faster (Section 4.2) and the network denser (Section 4.3), what happens when both decisions adjust (Section 4.4), how the design moves as the price rises (Section 4.5), and what changes under a cap on realized carbon (Section 4.6). Proofs are in Appendix A.

With $G(n, v) = \partial \mathrm{SS}(n, v)/\partial n$ denoting the marginal safety stock of an additional DC, differentiation gives the stationarity conditions

$$F_\tau + H_\tau G(n, v) = \frac{\ell_\tau N\mu\delta}{2}\, n^{-3/2}, \tag{11}$$

$$N\mu D\, T'_\tau(v) = \frac{H_\tau \,\mathrm{SS}(n, v)}{2v} + \frac{h_p\, P(v)}{v}, \tag{12}$$

the first balancing the marginal facility and lost-pooling costs against the marginal last-mile saving, and the second balancing the marginal transport cost of speed against the marginal holding and

financing saving it obtains. The speed condition (12) governs the *availability margin*: safety stock and replenishment speed are alternative ways to provide service. The density condition (11) governs the *proximity margin*: facilities and last-mile distance are alternative ways to reach markets.

4.1 Does the design respond to a carbon price predictably?

A firm that prices or caps its carbon needs a design that moves with the price in a predictable way. The objective is not jointly convex in n and v, but it is convex in their logarithms. Let $J^{\mathrm{r}}(n;\tau) = \min_{v_{\min}\le v\le v_{\max}} J(n,v;\tau)$ be the *reduced objective*, the cost, carbon charge included, of operating n distribution centers when speed is chosen optimally for that network. A function on an interval is *strictly quasiconvex* if it strictly decreases and then strictly increases, either phase possibly empty, so that it has a single minimizer.

Lemma 4.1 (A unique optimum)**.** *Let t and ε be increasing, differentiable, and convex, let all cost and carbon rates be positive, let $h_p \ge 0$, and let $v_{\min} > 0$. For every carbon price $\tau \ge 0$, and for the carbon objective C:* (i) *the objective is strictly convex in $(\ln n, \ln v)$ on the feasible region;* (ii) *the problem has a unique minimizer, and under a carbon price the minimizer $\big(n^*(\tau), v^*(\tau)\big)$ is continuous in τ;* (iii) *the speed response $v^*(n)$ is unique, continuous, and nondecreasing in n, and the density response $n^*(v)$ is nondecreasing in v, each strictly where it is interior when $\rho < 1$; the objective at any fixed speed and the reduced objective are strictly quasiconvex in n on $[1, N]$.*

Adding a distribution center shortens every delivery but adds a facility charge and the safety stock lost to pooling. Measured in proportional changes of the network and of speed, each of these effects is convex, so once an additional center stops paying for itself, no larger network makes it pay again. Because the objective is strictly convex in the logarithms of n and v, the stationarity conditions (11)–(12) identify the optimum whenever it is interior, and the comparisons below hold globally. The reduced objective is flat near its minimum, so treating the number of distribution centers as continuous costs little: in the calibrated network of Section 5, rounding it to the best integer raises the objective by at most 0.09%, in line with the accuracy reported for continuous-approximation designs (Ouyang and Daganzo 2006; Ansari et al. 2018). By part (iii), the two decisions are complements: faster replenishment shrinks safety stock and lowers the inventory penalty of dispersion, and a denser network raises the inventory at stake and the value of speed. A firm should therefore evaluate speed investments and network density together.

4.2 When should replenishment be faster?

Faster replenishment raises inbound transport cost and carbon but lowers the inventory the firm must hold. For a given network, the carbon-minimizing firm therefore compares the carbon per dollar of the inventory that speed saves with that of the speed itself.

Proposition 4.2 (Speed)**.** *Fix the density n, and let the speed-dependent cost and carbon be*

$$\hat{K}(v) = h\,\kappa(n)\,v^{-1/2} + h_p N\mu D\,v^{-1} + N\mu D\,t(v), \qquad \hat{C}(v) = \eta\,\kappa(n)\,v^{-1/2} + N\mu D\,\varepsilon(v),$$

with $\kappa(n) = z\sigma\sqrt{ND}\,\sqrt{\Psi(n)}$*, and let* t *and* ε *be increasing, differentiable, and convex.*

(i) *For interior optima, the carbon-minimizing speed exceeds the cost-minimizing speed if and only if the inventory saving from speed is more carbon-intensive per dollar than inbound transport* at the margin,

$$a(n, v_{\mathrm{K}})\,\frac{\eta}{h} \;>\; \frac{\varepsilon'(v_{\mathrm{K}})}{t'(v_{\mathrm{K}})}, \tag{13}$$

evaluated at the cost-minimizing speed v_{K}*. On* $[v_{\min}, v_{\max}]$ *the optima are the projections of the unconstrained speeds onto the interval, so the comparison holds weakly, and strictly when neither projection binds. With the quadratic instantiation* $\varepsilon(v) = \varepsilon_0 v^2$*,* $t(v) = t_0 + t_1 v^2$*, the right side of* (13) *is* r_t*; without financing (*$h_p = 0$*),* $a = 1$*, the threshold is* $r_h > r_t$*, and the unconstrained optimum is*

$$v^*(n) \;=\; \left[\frac{(h + \tau\eta)\, z\sigma\sqrt{ND}\,\sqrt{\Psi(n)}}{4\,(t_1 + \tau\varepsilon_0)\, N\mu D}\right]^{2/5}. \tag{14}$$

(ii) *Under the quadratic instantiation,* $v^*(n)$ *increases with the safety factor* $z = \Phi^{-1}(\alpha)$ *and with the demand standard deviation* σ *wherever it is interior. Without financing,* z *and* σ *leave* (13) *unchanged; with financing they raise* $a(n, v_{\mathrm{K}})$*, so they can turn the comparison toward faster replenishment but never toward slower.*

Total carbon is U-shaped in speed, so a firm that tightens its carbon target may need to replenish *faster*. By (3), a lower fuel share of the marginal cost of speed lowers the right side of (13), so speed bought with labor is more likely to pay in carbon than speed bought with fuel. Financing stock in transit works the other way: it makes speed valuable to the cost-minimizing firm for the capital it saves; capital carries no carbon, so financing lowers the share a. A higher service requirement or more volatile demand raises the speed at both optima and, with financing, the share a as well. The comparison is at a common density; when both decisions adjust, the density change reinforces the speed reversal if carbon also disperses the network and attenuates it otherwise (Section 4.4).

4.3 When should the network be denser?

Increasing density shortens the last mile but raises facility cost and carbon and weakens pooling. For a given speed, the carbon-minimizing firm therefore compares the carbon per dollar of last-mile delivery with a blend of facility and lost-pooling carbon.

Proposition 4.3 (Density)**.** *Fix the speed* v *and let*

$$G(n, v) \;=\; \frac{\partial\,\mathrm{SS}(n, v)}{\partial n} \;=\; z\sigma\sqrt{D/v}\,\sqrt{N}\,\frac{1-\rho}{2\sqrt{\Psi(n)}} \;\geq\; 0 \tag{15}$$

be the marginal safety-stock penalty of an additional DC.

(i) *The minimizer of* $J(\cdot, v)$ *is unique; it is interior if and only if* $J_n(1, v) < 0 < J_n(N, v)$*, and otherwise lies at a boundary,* $n^* \in \{1, N\}$*. For an interior cost-minimizing density* $n_K(v)$*, the*

carbon-minimizing density exceeds it, $n_C(v) > n_K(v)$, if and only if last-mile transport is more carbon-intensive per dollar than the pooling-adjusted facility threshold*,*

$$\frac{\varepsilon_L}{t_L} > \frac{\varphi + \eta\, G\big(n_K(v), v\big)}{F + h\, G\big(n_K(v), v\big)}. \tag{16}$$

The right-hand side is a weighted average of r_F and r_h, with weight on r_h rising in the pooling penalty G. In the no-pooling limit ($\rho \to 1$, so $G \to 0$) the optimum is closed form when interior, $n^* = (\ell_\tau N \mu \delta / 2F_\tau)^{2/3}$, and (16) *reduces to*

$$\frac{\varepsilon_L}{t_L} > \frac{\varphi}{F}. \tag{17}$$

(ii) *Under the quadratic instantiation, and with speed chosen optimally for each network, a higher z or σ raises the inventory penalty of a denser network and lowers the optimal density at every carbon price, strictly when it is interior; in particular, it shifts the carbon-minimizing density toward consolidation.*

The density decision is the centralize-versus-decentralize question in carbon terms. The carbon-minimizing network is more dispersed when a dollar of last-mile delivery emits more than the facility threshold, which pooling shifts from r_F toward r_h. Pooling therefore enters as a weight on the holding ratio: when holding is more carbon-intensive per dollar than facilities, stronger pooling raises the threshold for dispersion, and otherwise it lowers it. A corner optimum, when it arises, is fully centralized or fully decentralized. By part (ii), a higher service requirement or more volatile demand raises the inventory penalty of each additional center and so favors fewer distribution centers at every carbon price; endogenous service choice is a natural extension (Section 7).

4.4 What happens when both decisions adjust?

When both decisions adjust, complementarity links them: a carbon price that disperses the network also raises the value of speed, and one that speeds replenishment lowers the inventory penalty of dispersion. We use

$$R_{\text{speed}} = \frac{r_h}{r_t}, \qquad R_{\text{density}} = \frac{r_L}{r_F}$$

to denote the two diagnostic ratios, (n_K, v_K) and (n_C, v_C) the cost- and carbon-minimizing configurations, $\hat{v}_C$ the carbon-minimizing speed at the cost-minimizing density n_K, and $a_K = a(n_K, v_K)$ the safety-stock share at the cost minimum. Two quantities of the cost-minimizing network enter. The first is the *pooling weight*

$$\omega = \frac{h\, G(n_K, v_K)}{F + h\, G(n_K, v_K)} = 1 - \Big(\frac{n_K}{n_K^0}\Big)^{3/2}, \qquad n_K^0 = \Big(\frac{t_L N \mu \delta}{2F}\Big)^{2/3}, \tag{18}$$

the share of the marginal cost of a distribution center that is lost pooling. Since n_K^0 is the density the firm would choose if pooling were absent, ω measures how far pooling holds the network below

that size. The second, for interior n_K, is the pair of *speed bands*

$$\underline{R} = \sqrt{\Psi(n_K)/N} \;<\; 1 \;<\; \overline{R} = \sqrt{\Psi(n_K)/\Psi(1)}, \tag{19}$$

which bound how far a change in density can move the carbon-minimizing speed.

Proposition 4.4 (Joint design)**.** *Suppose that $\rho < 1$ and that the cost-minimizing density n_K is interior.*

(i) Density. *The carbon-minimizing network is more dispersed than the cost-minimizing one, $n_C > n_K$, if and only if*

$$r_L \;>\; (1-\omega)\, r_F \;+\; \omega\, r_h \sqrt{v_K/\hat{v}_C}, \tag{20}$$

with $n_C = n_K$ at equality and $n_C < n_K$ when the inequality is reversed. Under the quadratic instantiation, and when v_K and $\hat{v}_C$ are interior, (20) *becomes*

$$r_L \;>\; (1-\omega)\, r_F \;+\; \omega\, r_h^{4/5}\, r_t^{1/5}\, a_K^{-1/5}. \tag{21}$$

(ii) Speed. *Under the quadratic instantiation, and when the speed responses of the cost and the carbon objective are interior at n_K and at n_C, the carbon-minimizing speed satisfies*

$$\ln \frac{v_C}{v_K} = \frac{2}{5} \ln \left(a_K R_{\text{speed}}\right) + \frac{1}{5} \ln \frac{\Psi(n_C)}{\Psi(n_K)}, \tag{22}$$

so that $v_C > v_K$ if and only if $a_K R_{\text{speed}} > \sqrt{\Psi(n_K)/\Psi(n_C)}$. If (21) *holds and $a_K R_{\text{speed}} \geq 1$, the carbon-minimizing design is faster and more dispersed; if* (21) *is reversed and $a_K R_{\text{speed}} \leq 1$, it is slower and more consolidated. Whatever the density, it is faster if $a_K R_{\text{speed}} > \overline{R}$ and slower if $a_K R_{\text{speed}} < \underline{R}$; only between the bands, with the two conditions pointing in opposite directions, does the size of the density change decide.*

(iii) Magnitudes without pooling. *Under the quadratic instantiation, in the no-pooling limit ($\rho \to 1$), $\omega = 0$, the two decisions separate, and for interior optima*

$$\frac{n_C}{n_K} = R_{\text{density}}^{2/3}, \qquad \frac{v_C}{v_K} = \left(a_K R_{\text{speed}}\right)^{2/5}.$$

For a given product, Proposition 4.4 tells the firm whether to replenish faster or slower and whether to operate more or fewer distribution centers. The density condition (20) is Proposition 4.3 with one change. Adding a distribution center incurs a facility charge and lost pooling, in the proportions $1-\omega$ and ω of its marginal cost. The carbon-minimizing firm replenishes its centers at its own speed $\hat{v}_C$, and safety stock falls with the square root of speed, so per dollar of lost pooling its carbon is $r_h\sqrt{v_K/\hat{v}_C}$. The carbon-minimizing network is more dispersed when a dollar of last-mile delivery emits more than this blend of facility and lost-pooling carbon, whether or not speed is at a bound. When the speeds are interior, the carbon-minimizing firm limits the safety stock that an additional center adds by replenishing faster, so the center raises inbound freight carbon as well as

holding carbon (21). Unlike the marginal comparisons, the joint design therefore lets the carbon intensity of inbound freight bear on how many distribution centers to operate. The effect is modest: it is scaled by the pooling weight, freight carries only a small part of the carbon of lost pooling, and it vanishes where the carbon-minimizing speed stays at a bound. Financing stock in transit raises the cost-minimizing speed, and with it the carbon of lost pooling, so it raises the threshold for dispersion.

The speed decomposition (22) separates the two reasons a carbon-minimizing chain replenishes at a different speed. The intensity effect, $\frac{2}{5}\ln(a_K R_{\text{speed}})$, is the comparison of Proposition 4.2. The density effect, $\frac{1}{5}\ln[\Psi(n_C)/\Psi(n_K)]$, is complementarity: a more dispersed network holds more safety stock and so values speed more. The bands (19) bound the density effect, because no network pools more than full consolidation or less than full dispersion.

Part (iii) gives the size of the redesign. Without pooling and with both optima interior, a product whose last mile is four times as carbon-intensive per dollar as its facilities runs $4^{2/3} \approx 2.5$ times as many distribution centers when minimizing carbon. As we show in the proof of part (iii), the carbon that the cost-minimizing design leaves unabated is, to second order,

$$\frac{C(n_K, v_K)}{C_{\min}} - 1 \;\approx\; \theta\,\frac{(\ln R_{\text{density}})^2}{9} + (1-\theta)\,\frac{2\big[\ln(a_K R_{\text{speed}})\big]^2}{25}, \tag{23}$$

where $C_{\min}$ is the carbon minimum and θ is the share of proximity carbon, $\varphi n + \varepsilon_L N \mu d(n)$, in that minimum. The excess is small near parity and rises quickly away from it: the cost-minimizing network emits about 5% more proximity carbon than the minimum when $R_{\text{density}} = 2$ and about 19% more when $R_{\text{density}} = 4$. Carbon-aware redesign therefore matters most for products whose ratios lie far from one. With pooling the two decisions interact and these expressions are no longer exact; Section 5 illustrates the magnitudes in that case.

4.5 How does the design move as the carbon price rises?

A firm facing a rising carbon price also needs to know how the design moves along the way. For an activity with carbon intensity per dollar r_x, we use $s_x(\tau) = \tau r_x/(1+\tau r_x)$ to describe the *carbon share* of its blended rate — for holding, $s_h = \tau\eta/(h+\tau\eta)$. The share rises from zero to one as the price rises.

Proposition 4.5 (Carbon-price path)**.** *Under the quadratic instantiation:* (i) *at a fixed density and an interior speed response,* $\mathrm{d}\ln v^*/\mathrm{d}\ln\tau = (a_\tau s_h - s_t)/(3 - a_\tau/2)$*, where* $a_\tau = H_\tau \mathrm{SS}/(H_\tau \mathrm{SS} + 2h_p P)$ *at the optimum; its sign is that of* $a(n, v^*)\,r_h - r_t$*, so the speed moves monotonically in* τ*, and without financing the elasticity is* $\frac{2}{5}(s_h - s_t)$*;* (ii) *in the no-pooling limit and for an interior density,* $\mathrm{d}\ln n^*/\mathrm{d}\ln\tau = \frac{2}{3}(s_L - s_F)$*;* (iii) *without financing, with pooling and interior speeds, if* r_L *exceeds each of* r_F*,* r_h*, and* r_t*, the optimal density* $n^*(\tau)$ *is nondecreasing along the whole price path, and if* r_L *falls short of each, it is nonincreasing.*

The carbon share of a blended rate measures how much of that activity's cost, carbon charge

included, is carbon. By part (i), for a given network, a rising carbon price speeds replenishment exactly when the inventory saving from speed is more carbon-intensive per dollar than inbound freight, and the speed moves in one direction along the whole path. The response is gradual: at a low price both shares are near zero and speed barely moves, while at a high price both approach one and speed settles at its carbon-minimizing level. The speed moves most where the carbon shares of holding and inbound freight differ most, as when the speed the firm buys is fuel-intensive and its held stock carries little carbon. Part (iii) gives a simple condition for density. Without financing, when the last mile is the most carbon-intensive of the four activities per dollar and speeds are interior, no increase in the carbon price calls for consolidation; when the last mile is the least carbon-intensive, no increase calls for dispersion.

4.6 What changes under a cap on realized carbon?

A firm held to a carbon budget for each accounting period must meet it with the carbon it actually emits, which fluctuates with demand about $C(n, v)$: the facility charge and the base-stock safety-stock target are planned and treated as deterministic, but inbound and last-mile freight scale with realized throughput $Q = \sum_i X_i$, the sum of the realized market demands, which has mean $N\mu$ and variance $\sigma_Q^2 = \sigma^2 N\,[\,1 + (N-1)\rho\,]$. Realized carbon is then

$$\tilde{C}(n, v) \;=\; \varphi n + \eta\,\mathrm{SS}(n, v) + \big[\varepsilon(v)D + \varepsilon_L d(n)\big]\,Q, \tag{24}$$

and a firm that requires realized carbon to stay within B with confidence $1-\beta_c$ imposes the chance constraint $\Pr\{\tilde{C}(n, v) \le B\} \ge 1 - \beta_c$, the carbon-side counterpart of the cycle service level α on the inventory side.

Corollary 4.6 (Caps on realized carbon)**.** (i) *The chance constraint is equivalent to*

$$\varphi n + \eta\,\mathrm{SS}(n, v) + m\big[\varepsilon(v)N\mu D + \varepsilon_L N\mu\, d(n)\big] \;\le\; B, \tag{25}$$

the carbon function (7) *with the transport component scaled by* $m = q_{1-\beta_c}/(N\mu)$*, where* $q_{1-\beta_c}$ *is the lower* $(1-\beta_c)$ *quantile of realized throughput. Under the normal approximation,* $m = 1 + z_{\beta_c}\,\mathrm{CV}_Q$*, with* $z_{\beta_c} = \Phi^{-1}(1-\beta_c)$ *and* $\mathrm{CV}_Q = \sigma_Q/(N\mu)$*, so* $m \ge 1$ *whenever* $\beta_c \le \frac{1}{2}$*, and, for* $\beta_c < \frac{1}{2}$*,* m *rises with the demand coefficient of variation* σ/μ*, the correlation* ρ*, and the confidence* $1-\beta_c$*.*

(ii) *The speed threshold* (13) *becomes* $a(n, v_{\mathrm{K}})\,\eta/h > m\,\varepsilon'(v_{\mathrm{K}})/t'(v_{\mathrm{K}})$*, the density threshold* (16) *becomes* $m\,\varepsilon_L/t_L > (\varphi+\eta G)/(F+hG)$*, and in the joint design the speed ratio becomes* $a_K R_{\mathrm{speed}}/m$ *and* (21) *becomes* $m\,r_L > (1-\omega)\,r_F + \omega\,m^{1/5} r_h^{4/5} r_t^{1/5} a_K^{-1/5}$*. Under the quadratic instantiation, the carbon-minimizing speed is nonincreasing in* m*, and strictly decreasing where* n_C *and* v_C *are interior.*

A cap on realized carbon therefore leaves the structure intact. It places a risk margin on transport emissions, just as safety stock places one on demand, and so reinforces slower replenishment and, where the last mile is carbon-intensive, more distribution centers; the expected-value cap is the

limiting case $m = 1$. The deterministic equivalent is distribution-free in the throughput quantile, and a firm with strongly skewed demand can replace the normal approximation by the empirical quantile. The cap applies over the same accounting period as the flow model — here a week. For a cap on cumulative carbon over T periods with demands independent across periods, throughput has coefficient of variation $\mathrm{CV}_Q/\sqrt{T}$, so the margin shrinks as the horizon lengthens: in the calibrated network of Section 5, the 95% multiplier falls from $m = 1.37$ for a weekly cap to $m = 1.05$ for an annual cap ($T = 52$). The margin therefore matters most for high-frequency operational carbon budgets and less for annual compliance accounting.

5 A Calibrated Application

This section applies the results to distribution activities measured from public US data. Section 5.1 constructs benchmark carbon intensities per dollar from public sources and stated assumptions, Section 5.2 traces how a reference network meets a carbon target, and Section 5.3 examines whether that design survives changes in the assumptions. The directions follow from the propositions once the ratios are known, subject to the bounds on speed; the magnitudes describe the network we calibrate.

5.1 Benchmark carbon intensities

Each ratio is the carbon an activity emits per dollar it costs. We compute the transport ratios from the bottom up, as the activity's energy use times the emission factor of that energy, divided by the activity's cost, and compare the result with the industry factors of the US environmentally extended input–output model (U.S. Environmental Protection Agency 2024a); the warehouse ratios come from that model. A dollar spent on energy carries about 2.8 kg CO_2e as diesel, 2.9–4.6 kg as grid electricity, and 5.0–5.8 kg as marine fuel, while the labor and capital that make up the rest of logistics cost carry little carbon of their own. An activity's carbon intensity per dollar therefore rises with the share of its cost spent on energy.

For warehouses, we use the EPA input–output rate for warehousing and storage, 0.244 kg CO_2e per dollar (U.S. Environmental Protection Agency 2024a). It applies to the fixed operating cost of a distribution center, which gives the facility carbon φ, and to the storage charge that varies with the stock held, which gives η; the storage charge of a third-party warehouse is split by a share γ that is variable for an owned facility, and we take $\gamma = \frac{1}{2}$. The rate does not distinguish refrigerated from ambient storage, so refrigeration enters through the higher cost of the center. The rate counts the emissions of the inputs a warehouse buys as well as those of its own energy, and it lies near the upper end of what public data support: Lineage reports 0.28 kg per dollar of warehousing revenue, and the site energy of a median ambient distribution center, divided equally between the building and the stock it holds, gives about a fifth of the rate. We report both alternatives in Section 5.3. Storage emissions per pallet depend on the storage environment and not on the value of the goods.

Table 2: Benchmark carbon intensity per dollar of distribution activities, kg CO_2e per US dollar, with ranges over the stated inputs and allocation shares.

Ratio	Activity	Central	Range
Holding, r_h	Ambient packaged goods	0.108	0.011–0.22
	High-value goods	0.006	0.000–0.031
	Refrigerated dairy	0.160	0.065–0.25
	Refrigerated dairy, spoilage 0.5% per week	0.463	0.33–0.56
Line-haul speed, r_t	Faster vessel speed	5.37	4.97–5.82
	Truck instead of intermodal rail	1.78	1.18–3.55
	Air instead of truck	1.10	0.60–1.90
	Team instead of solo drivers	0.03	0.00–0.10
Last mile, r_L	Diesel local truck	0.50	0.30–0.60
	Diesel van	0.30	0.15–0.45
	Electric truck	0.16	0.12–0.20
	Electric van	0.075	0.05–0.10
Facilities, r_F	Ambient distribution center	0.244	0.015–0.28
	Refrigerated distribution center	0.244	0.027–0.28

Table 2 reports the ratios.[2]

Without spoilage, holding carries less carbon per dollar than a distribution center, because its cost is mostly the capital tied up in stock and the space and labor to store it, and the ratio falls with value density. Speed bought with fuel — a faster vessel, truck instead of intermodal rail, air instead of truck — is the most carbon-intensive, because the added cost is mostly fuel. A team of two drivers runs the same truck about twice as far per day, so the speed it buys adds mostly wages and little carbon. Spoilage changes the holding ratio of perishables: food that spoils while held carries the carbon of its production, which we value at the input–output intensity of the food. Even beef, the highest of the food categories we examine at 1.12 kg per dollar of processed product,

[2]Sourced inputs. Emission factors: diesel 10.21 kg CO_2 per gallon and the US grid 0.371 kg CO_2e per kWh (U.S. Environmental Protection Agency 2022); marine fuel 3.114 t CO_2 per metric ton (European Commission 2016). Prices: diesel \$3.66 per gallon (2025 average) (U.S. Energy Information Administration 2026), electricity 8.1–12.8 cents per kWh (U.S. Energy Information Administration 2025), and marine fuel at average prices of \$535 (2025) and \$626 (2024) per metric ton (Ship & Bunker 2026). Warehouses: the input–output rate for warehousing and storage (U.S. Environmental Protection Agency 2024a); for the ranges and the site-energy alternative, site energy of 22.7 (ambient) and 84.1 (refrigerated) kBtu per square foot (U.S. Environmental Protection Agency 2024b), about 30 (ambient) and 109 (refrigerated) kg CO_2e per pallet position-year, and 0.28 kg per dollar of warehousing revenue (Lineage, Inc. 2024); industrial rent of \$10.12 per square foot (Cushman & Wakefield 2025). Line haul: truckload at \$2.26 per mile, of which fuel is \$0.48 (American Transportation Research Institute 2025). Last mile: the courier and local-trucking industry factors (U.S. Environmental Protection Agency 2024a) and electric-van energy use of 0.37–0.39 kWh per mile (Mulholland 2022). Food: input–output intensities of processed beef, fluid milk, and poultry (U.S. Environmental Protection Agency 2024a). Assumptions. Third-party storage at \$240 (ambient) and \$480 (refrigerated) per pallet-year; rent at about 40% of a distribution center's fixed operating cost, with cold-storage rent twice ambient; capital at 10% of pallet value, for stock held and stock in transit; intermodal 10–30% cheaper than truckload with about 60% less carbon per load; a team premium of about \$0.35 per mile, carrying about 0.03 kg CO_2e per dollar, at twice the daily distance of a solo driver; air at four to five times the cost of road; distribution centers of 120,000 square feet; truckloads of 26 pallets; pallet values; the storage cost share γ; and, for the site-energy alternative, the share β of building energy that varies with the stock held, $\frac{1}{2}$ ($\frac{1}{4}$ to $\frac{3}{4}$ for the ranges). Ranges combine the low and high values of every input, with warehouse carbon from site energy at the low end and Lineage's reported intensity at the high end.

Table 3: Inputs of the reference network.

Input	Model	Reference network
Markets and demand	N; μ, σ; ρ	40 markets; 250 pallets a week each, coefficient of variation 0.40; correlation 0.30
Service	z	95% cycle service ($z = 1.645$)
Distances	D; δ	1,000 miles inbound; $d(n) = 700\,n^{-1/2}$ miles to the markets
Line haul	v; $t(v)$; $\varepsilon(v)$	167 (carload rail) to 500 (truckload) miles a day, through intermodal at 250; \$0.066 to \$0.087 and 0.015 to 0.052 kg per pallet-mile
Store delivery	t_L; ε_L	diesel truck, \$2.75 a mile with 20 pallets a load; $r_L = 0.50$
Distribution center	F; φ	\$3.0 million and 732 t a year
Holding	h; η	\$5,120 and 29 kg per pallet-year
Stock in transit	h_p	\$5,000 per pallet-year, 10% of the value of the goods

stays below the ratio of buying speed with truck instead of intermodal rail or with vessel speed; air freight, whose ratio spans 0.60 to 1.90, is the one fuel-bought option against which spoilage could favor speed. The spoilage rate of 0.5% of stock value per week is an assumption.

5.2 Meeting a carbon target in a reference network

We calibrate the model to a distributor that replenishes stores in forty metropolitan markets with ambient goods worth \$50,000 a pallet and that owns the goods while they travel. Table 3 maps each input of the model to its counterpart in this network. Door to door on a 1,000-mile lane, truckload takes about two days, since a solo driver covers about 500 miles a day (Federal Railroad Administration 2013); intermodal rail takes about four, of which three are from ramp to ramp and the rest local trucking at each end; and carload rail takes six days or more (Kwon et al. 1995). We fit $t(v)$ and $\varepsilon(v)$ through intermodal and truckload and extend them to a six-day transit, so the firm can replenish at any speed between 167 and 500 miles a day. At six days the fitted carbon per pallet-mile is 29% of truckload's, close to the fuel efficiency of carload moves relative to truck on competing corridors (ICF International 2009). The fitted carbon rate includes a constant term, the carbon per pallet-mile that does not vary with speed; it leaves the design at a given carbon price unchanged, but it enters total carbon and so the percentage cuts and the caps they define.

At the cost minimum, with the number of centers treated as continuous, the network operates 3.9 distribution centers and ships in 2.15 days, close to truckload, because financing makes speed valuable: the stock in transit, 3,066 pallets, exceeds the safety stock, 2,210, so the safety-stock share is $a_K = 0.27$. The pooling weight is $\omega = 0.08$. Holding emits 0.006 kg per dollar and the speed that truckload buys over intermodal rail 1.78, so the carbon pressure on speed points toward slower replenishment (Proposition 4.2). Diesel delivery emits 0.50 kg per dollar, above the threshold of 0.22 in Proposition 4.4(i), so the carbon pressure on density points toward more distribution centers.

Two features of the reference network account for the size of the cuts that follow. First, the slower end of the speed range corresponds to rail services, which emit 60 to 71% less per load than

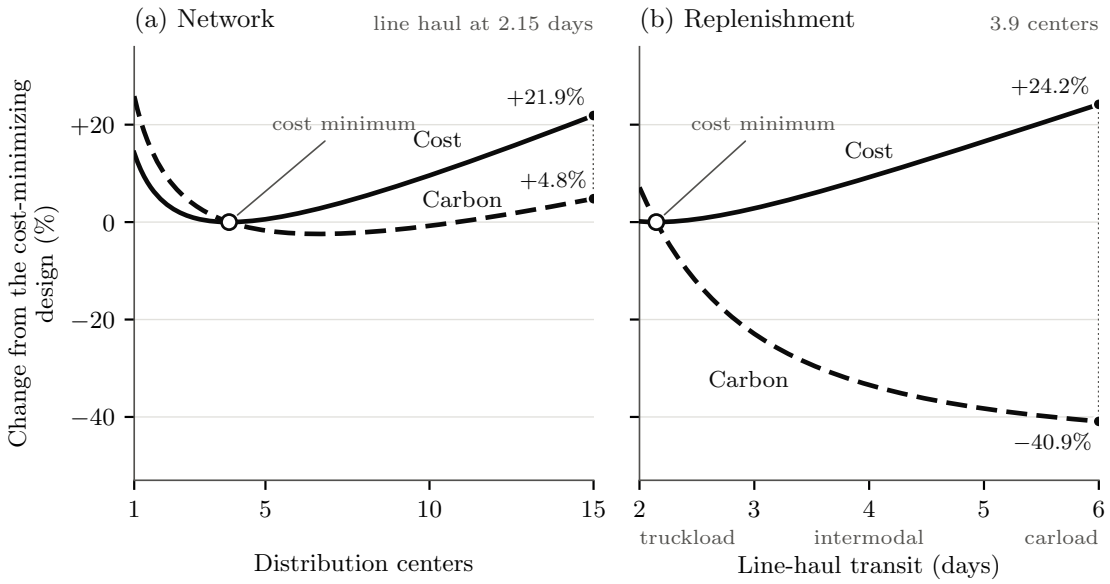


Figure 2: Annual cost and carbon of the reference network when one decision moves and the other stays at the cost-minimizing design, which costs $107.5 million and emits 39,659 t a year.

truckload, so the fitted carbon rate falls steeply as the line haul slows. Second, transport accounts for 93% of the network's logistics carbon, 61% on the inbound lane and 32% on store delivery, and warehousing for 7%. Studies of the logistics sector put warehousing at 11 to 20% (World Bank 2021; McKinsey & Company 2024); in the reference network the inbound lane and the delivery legs carry more of the total.

Figure 2 shows how annual cost and carbon change when one decision moves away from the cost-minimizing design and the other is held there. More distribution centers shorten the last-mile legs but add facility cost and carbon and weaken pooling. Carbon is lowest at 6.6 centers, where a further center adds as much facility and holding carbon as it saves on the last mile, and it is then only 2.4% lower, at a cost 2.8% higher; beyond that point more centers raise both cost and carbon, by 21.9 and 4.8% at fifteen centers. Slowing the line haul lowers freight cost and carbon but adds safety stock and stock in transit, which raise cost and add little carbon. At intermodal's four days, carbon is 33.4% lower and cost 9.2% higher; at six days, 40.9% lower and 24.2% higher. We next ask how the firm combines the two decisions to meet a required cut at least cost.

Figure 3 traces the least-cost design as the required cut in emissions grows, for three sets of actions: changing replenishment alone, changing the number of distribution centers alone, and changing both. Slower line haul delivers nearly all of each cut. At a 20% cut, the least-cost design slows the line haul from 2.15 to 2.79 days and grows the network from 3.9 to 4.3 centers, at a cost increase of 1.9%; lower inbound freight carbon accounts for 96% of the reduction. Changing both decisions is cheaper than slowing alone, and the advantage grows with the cut: a 30% cut raises cost by 5.8% against 6.3% by slowing alone, and a 40% cut by 17.0% against 21.0%. Changing both decisions attains at most a 43.3% cut, the carbon minimum with 6.6 centers, against the

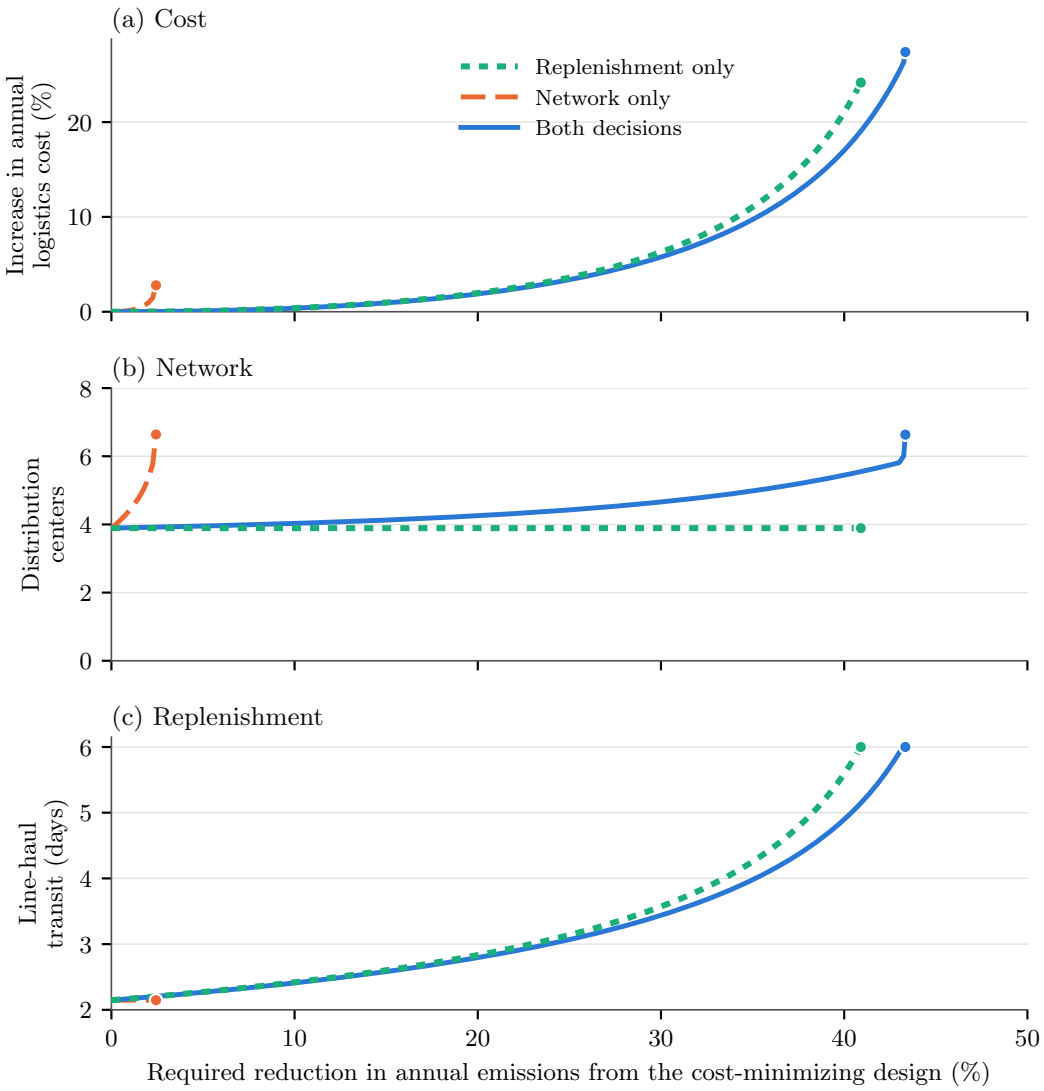


Figure 3: Least-cost design of the reference network for each required cut in emissions, for three sets of actions. Each curve ends at the largest cut that its actions attain.

40.9 and 2.4% that each decision attains alone (Figure 2). Each cut corresponds to a carbon price (Lemma 4.1). A 30% cut costs \$521 per metric ton avoided on average and corresponds to a price of \$1,632 per metric ton. A price of \$100 per metric ton lowers emissions by 5.2% at a cost increase of 0.1%, and the average metric ton avoided costs \$48, less than the price.

In Figure 4, we vary r_t at a fixed cost of speed. Panel (a) shows the mechanism of Proposition 4.4 for perishables that spoil, on a lane served by solo or team drivers in two days or one. An additional center raises safety stock, and faster replenishment wins it back. When speed is bought with labor, as with team drivers, winning it back costs little carbon: with the higher pooling weight, $\omega = 0.10$ (coefficient of variation 0.60, correlation 0.1, 99% service, and centers at \$1.0 million a year), the carbon-minimizing design ships in one day with 12.8 centers. As r_t rises, the line haul slows to

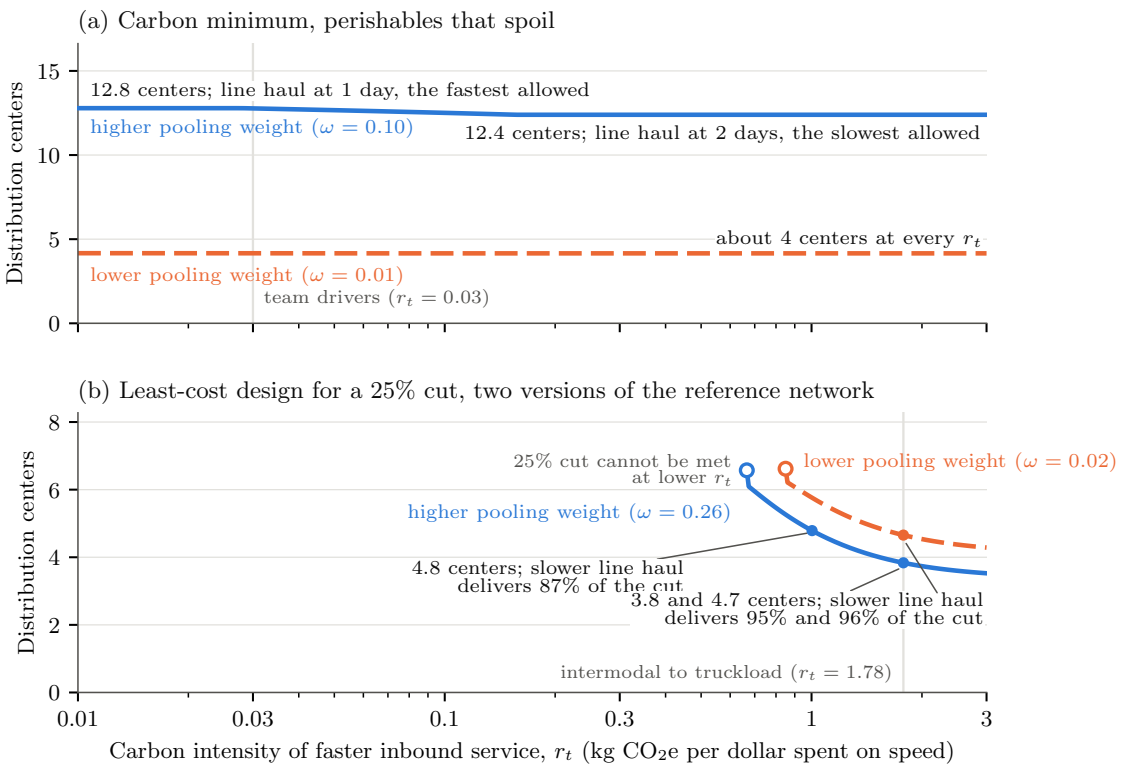


Figure 4: Number of distribution centers against the carbon intensity of faster inbound service, with delivery by diesel truck. Open circles mark the lowest r_t at which the 25% cut can be met.

two days and the network falls to 12.4 centers, and from there r_t no longer matters. With the lower pooling weight, $\omega = 0.01$ (the reference demand and service, and centers at \$6.0 million a year), density does not respond to r_t. Panel (b) shows a second channel: the least-cost design for a 25% cut in two versions of the reference network, one with the demand and service of panel (a) ($\omega = 0.26$) and one with coefficient of variation 0.20 and correlation 0.6 ($\omega = 0.02$). When speed is bought with fuel, as with truckload in place of intermodal, slower line haul delivers 95 to 96% of the cut and about four or five centers suffice. The less carbon faster service carries, the less slowing saves, and the more of the cut the network must deliver: at $r_t = 1.0$, under the higher pooling weight, slower line haul delivers 87% of the cut, with 4.8 centers. Because an added center emits nearly as much at the facility as it saves on the last mile, the network can deliver little of the cut, and below r_t of about 0.7, or 0.85 under the lower pooling weight, the cut cannot be met. Each line ends where only the carbon-minimizing network meets the cut.

5.3 Robustness

Table 4 examines whether the least-cost design for a 30% cut survives changes in the assumptions. At that cut the reference network operates 4.7 centers and ships in 3.44 days, at a cost increase of 5.8%. After the reference row, two rows are implementation checks, three are parameter and

Table 4: Robustness of the least-cost design for a 30% cut in emissions. Days are line-haul transit times.

	Cost minimum		Least cost for a 30% cut			
	Centers	Days	Centers	Days	Emissions cut	Cost increase
Reference	3.9	2.15	4.7	3.44	30.0%	5.8%
Implementation checks						
Truckload, intermodal, or carload only	3.9	2.00	3.8	4.00	37.7%	9.0%
Whole number of centers	4	2.15	5	3.42	30.0%	5.9%
Parameter and accounting sensitivities						
No financing of stock in transit	3.8	3.63	not attainable; largest cut 18.6%			
Annual cap on realized carbon, 95%	3.9	2.15	4.9	3.72	33.0%	7.9%
Weekly cap on realized carbon, 95%	3.9	2.15	not attainable			
Model extension						
Load-sensitive delivery, $\lambda = 0.25$	2.2	2.15	2.8	3.54	30.0%	6.4%

accounting sensitivities, and the last is a model extension. Each row reports the cost-minimizing design and the least-cost design for a 30% cut from that row's own cost minimum. For the caps on realized carbon, the budget is that of the reference row, and demand is independent across weeks.

In every attainable row the cut is met by slowing the line haul and, except when speed is restricted to the three services, by adding at most about one center. With only the three services the firm moves from truckload to intermodal, which cuts emissions by 37.7% at a cost increase of 9.0% with 3.8 centers in place of 3.9. In the reference network, the services available on the lane thus change the network recommendation, and the designs of Figure 3 depend on the fitted rates between and beyond the observed services. With whole centers, the design for a 30% cut has five centers and a line haul of 3.42 days, at a cost increase of 5.9%, close to the continuous design. A carbon price, however, then selects designs that jump between cuts: at about \$1,140 per metric ton the priced design moves from four centers and a 25% cut to five centers and a 27% cut, so a cut between the two is met at least cost by a design that no price selects. The gap is small, in line with the approximate equivalence of a carbon tax and a cap that Comas Martí et al. (2015) report for a discrete network. Without financing, the cost-minimizing chain, with 3.8 centers, already ships at 3.63 days, so slowing saves less and the largest attainable cut is 18.6%. A cap on realized carbon at the reference budget requires a deeper expected cut: 33.0% with 4.9 centers and 3.72 days, at a cost increase of 7.9%, for an annual cap, while a weekly cap at that budget cannot be met. In the extension, last-mile cost and carbon per pallet scale with $(n/4)^{\lambda}$, so that vehicle fill falls as the network disperses; with $\lambda = 0.25$ the cost-minimizing network has 2.2 centers, and the cut raises it to 2.8 and slows the line haul to 3.54 days, at a cost increase of 6.4%.

Table 5 varies the carbon inputs of the reference network and its storage environment. Across the rows, warehousing accounts for 1.6 to 9.4% of carbon, and a 30% cut needs 3.1 to 5.6 centers, costs 3.0 to 17.1% more than the cost minimum, or \$328 to \$1,516 per metric ton avoided, and corresponds to a carbon price of \$923 to \$7,357 per metric ton. The carbon saving of intermodal

Table 5: Sensitivity of the least-cost design for a 30% cut in emissions to the carbon inputs and the storage environment.

	Warehousing	Least cost for a 30% cut			
	share of carbon	Centers	Cost increase	Cost per metric ton avoided	Carbon price
Reference	7.3%	4.7	5.8%	\$521	\$1,632
Warehouse carbon from site energy	1.6%	5.0	4.2%	\$401	\$1,110
Warehouse carbon at 0.28 kg per dollar	8.3%	4.6	6.1%	\$545	\$1,747
Truckload emission factor 20% lower	8.4%	4.9	7.5%	\$775	\$2,612
Truckload emission factor 20% higher	6.6%	4.5	4.8%	\$385	\$1,156
Intermodal 45% below truckload	7.2%	5.6	17.1%	\$1,516	\$7,357
Intermodal 70% below truckload	7.4%	4.4	3.6%	\$328	\$923
Store delivery, $r_L = 0.30$	8.4%	4.0	4.1%	\$423	\$1,243
Store delivery, $r_L = 0.60$	6.9%	5.0	6.6%	\$557	\$1,780
Electric store delivery, $r_L = 0.16$	9.4%	3.7	3.0%	\$345	\$951
Refrigerated storage	8.7%	3.1	7.2%	\$641	\$2,193

over truckload matters most. At 45% in place of 60%, slowing saves less carbon per day, and the cut needs 5.6 centers and costs 17.1%, or \$1,516 per metric ton, at a price of \$7,357; at 70%, it needs 4.4 centers and costs 3.6%, or \$328 per metric ton, at a price of \$923. A truckload emission factor 20% lower reduces the carbon that slowing saves, and the cut needs 4.9 centers and costs 7.5% (\$775 per metric ton, at a price of \$2,612); 20% higher, it needs 4.5 centers and costs 4.8% (\$385 per metric ton, at \$1,156). Warehouse carbon matters less, because transport dominates the network's emissions: from site energy alone, warehousing is 1.6% of carbon and the cut needs 5.0 centers and costs 4.2% (\$401 per metric ton, at \$1,110), and at 0.28 kg per dollar, the intensity Lineage reports for its warehouses, it is 8.3% and the cut needs 4.6 centers and costs 6.1% (\$545 per metric ton, at \$1,747). Store delivery at 0.30 and 0.60 kg per dollar calls for 4.0 and 5.0 centers, at cost increases of 4.1 and 6.6% (\$423 and \$557 per metric ton, at \$1,243 and \$1,780). Electric delivery, whose carbon per dollar lies below that of a distribution center, calls for fewer centers, 3.7, as Proposition 4.4 predicts; warehousing is then 9.4% of carbon, and the cut costs 3.0% (\$345 per metric ton, at \$951). Refrigerated storage doubles the fixed cost of a center, so the network uses fewer centers, 3.1 for the cut, at a cost increase of 7.2% (\$641 per metric ton, at \$2,193).

6 Managerial Implications

The results give a firm a diagnostic it can compute from its operating rates, demand characteristics, and cost-minimizing configuration, and they indicate which of its distribution decisions a carbon price will move, in which direction, and by how much. Each implication below follows from a result of Section 4, and the calibrated values of Section 5 illustrate it.

A diagnostic in two questions. A firm needs four ratios and two shares. The ratios are the carbon per dollar of holding stock, of the speed it can buy on its line-haul lane, of last-mile delivery,

Table 6: Carbon-minimizing direction of the two decisions (Proposition 4.4), with examples from the central ratios of Table 2.

	More distribution centers *last mile more carbon-intensive per dollar than facilities and lost pooling*	**Fewer distribution centers** *last mile less carbon-intensive per dollar than facilities and lost pooling*
Faster replenishment *inventory saving from speed more carbon-intensive per dollar than the speed it buys*	team drivers for perishables that spoil; diesel delivery	team drivers for perishables that spoil; electric delivery
Slower replenishment *inventory saving from speed less carbon-intensive per dollar than the speed it buys*	speed bought with fuel, or team drivers for high-value goods; diesel delivery	speed bought with fuel, or team drivers for high-value goods; electric delivery

and of operating a distribution center. Labor and capital carry little carbon, so each transport ratio follows mainly from the activity's energy spend, the warehouse ratios from the input–output rate for warehousing and, for perishables, the production carbon of stock that spoils (Section 5.1). A firm can therefore estimate the ratios from its fuel, electricity, warehousing, and spoilage records. The shares are the pooling weight ω — the share of a distribution center's marginal cost that is lost pooling — and the safety-stock share a_K of the inventory that speed saves, and the firm can obtain both from its cost-minimizing network through (18) and (8). Proposition 4.4 then gives the direction of both decisions, and Table 6 sets out the four combinations; off the diagonal, the speed direction shown holds whenever $a_K R_{\text{speed}}$ lies outside the bands (19).

Slow the speed that is bought with fuel. A faster vessel, truck instead of intermodal rail, and air instead of truck add cost that is mostly fuel, so their carbon per dollar is many times that of holding stock (Table 2). By Proposition 4.2, a carbon price then favors slower replenishment, and financing the goods in transit strengthens that direction because it lowers the share a_K (Section 4.2). In the reference network, which finances its stock in transit, slower line haul delivers nearly all of a 20% cut in emissions (Figure 3). The shifts of pharmaceutical freight from air to ocean and of truckload lanes to intermodal rail are consistent with this direction (Pharmaceutical Commerce 2023; FreightWaves 2026), as is the finding of Maloni et al. (2013) that slower vessel speed lowers carbon and total cost despite the added pipeline inventory. Where the cost-minimizing chain already ships at the slowest option, a carbon price leaves speed there.

Buy speed with labor where held stock carries carbon. A team of drivers on the same truck delivers sooner for nearly the same carbon: the truck covers about twice the distance per day on the same fuel per mile, and the added cost is the second driver's wages. The shorter lead time lowers both the safety stock and the stock in transit. Because the speed a team buys costs wages and little fuel, its carbon per dollar is close to zero, and a carbon price favors it wherever the inventory

saving from speed is more carbon-intensive per dollar than that speed (Proposition 4.2). That requires held stock to carry carbon. At the input–output rate for warehousing, holding ambient packaged goods emits 0.11 kg per dollar and holding perishables that spoil 0.46, about four and fifteen times the ratio of team operation (Table 2). For packaged goods and perishables that do not spoil, however, the cost- and carbon-minimizing speeds on our 1,000-mile lane both lie at solo operation, so a carbon price leaves speed there; in our calibration only perishables that spoil move to team drivers. For high-value goods, whose holding cost is mostly capital, team operation is the more carbon-intensive of the two, and a carbon price slows it. A team shortens lead time only where a trip exceeds the distance one driver may cover in a day, so speed bought with labor matters most on long lanes. Demand volatility and the service target raise the cost-minimizing and the carbon-minimizing speeds together; without financing they leave the direction unchanged, and with financing they raise a_K and can turn it toward faster replenishment, never slower (Proposition 4.2).

Reassess network density when delivery technology changes. The density decision compares the carbon per dollar of last-mile delivery with that of a distribution center's facility and lost pooling (Propositions 4.3 and 4.4). Diesel delivery, by truck or van, is carbon-intensive enough per dollar to call for more distribution centers closer to demand. At the central ratios of Table 2, an electric truck or van emits less per dollar than a distribution center, so a carbon price calls for fewer, larger centers, and in the reference network electric delivery lowers the number of centers for a 30% cut from 4.7 to 3.7 (Table 5). The direction of electric delivery depends on the warehouse rate: measured by site energy alone, a center emits less per dollar than an electric truck, and an ambient center less than an electric van. Amazon reports that placing inventory closer to customers avoided 127 million vehicle miles in 2025 (Amazon.com, Inc. 2026), a move in the direction Proposition 4.4 gives when the last mile is carbon-intensive. A firm that electrifies its last mile should therefore revisit how many distribution centers it operates. The carbon intensity of the line haul also bears on density, through the share of a carbon target that slowing can deliver and through the speed at which the firm replenishes its centers; by Proposition 4.4, the channel through speed is modest (Figure 4(a)).

Expect modest moves at current carbon prices. Proposition 4.5 ties the response to a carbon price to the carbon share of each activity's cost, carbon charge included. At \$100 per metric ton, carbon is about 1% of the cost of holding ambient goods, 2% of the cost of operating facilities, 5% of the cost of diesel truck delivery, 15% of speed bought with truck instead of intermodal rail, and 35% of speed bought with vessel speed; it reaches half of the cost of diesel truck delivery only at about \$2,000 per metric ton. In the reference network a price of \$100 per metric ton lowers emissions by 5.2% at a cost increase of 0.1%, almost entirely through slower line haul, and even a 30% cut, which corresponds to a price of \$1,632 per metric ton, grows the network by less than one center (Figure 3). The two decisions are complements (Lemma 4.1), so a program that acts on one alone, such as a slow-steaming mandate or a consolidation drive, evaluates it at a level of the other decision that the joint design would change.

Account for realized carbon over short horizons. A cap on realized rather than expected carbon places a risk margin on transport emissions (Corollary 4.6). For a weekly cap held with 95% confidence, the margin raises the transport carbon rates by 37% in our calibration and so reinforces slower fuel-bought speed and, with diesel delivery, more distribution centers. For a 30% cut in the reference network, a weekly cap at the expected budget cannot be met (Table 4). For an annual cap the margin is about 5%, and the design for the cut moves only to 4.9 centers and 3.72 days, at a cost increase of 7.9%.

Implications for policy. An instrument that prices or caps transport alone, such as a freight fuel levy or a vehicle emissions standard, blends the carbon price into the transport rates and leaves holding and facility carbon unpriced. At a given density it therefore slows replenishment wherever speed is above its lower bound, and at a given speed it adds distribution centers, whatever the product. By the marginal comparisons of Propositions 4.2 and 4.3, these moves lower total carbon, at least for small prices, when they agree with the carbon-minimizing direction, and raise it when they do not. At a given density, a transport-only price therefore raises total carbon when the inventory saving from speed is more carbon-intensive per dollar than the speed it would buy; at a given speed, it does so when the last mile is less carbon-intensive per dollar than facilities and lost pooling. When both decisions adjust, its effect on total carbon depends on their joint response. By the ratios of Table 2, the first condition holds for perishables that spoil served by team drivers, where the chain buys such speed, and the second for electric delivery. An instrument that prices total logistics carbon, including warehousing, lets the firm weigh the carbon of all four activities when it chooses speed and density.

7 Discussion and Conclusion

A firm that distributes through regional distribution centers provides availability through safety stock or replenishment speed and proximity through facilities or last-mile travel, and a carbon price or cap can change which activity provides each. We have shown that the joint design has a unique optimum at every carbon price and that the carbon-minimizing direction of density, and of speed when the speed and density comparisons agree, follows from four carbon-intensity-per-dollar ratios and two shares of the cost-minimizing design, the pooling weight and the safety-stock share. A firm that finances its goods in transit buys speed for capital that carries no carbon; financing raises the cost-minimizing speed without changing the carbon-minimizing one. In public US data, the classification depends on the product, on how speed is bought, and on how the last mile is driven. At the central ratios, carbon calls for slower fuel-bought speed for every product, and speed bought with labor is favored where held stock carries carbon, as for perishables that spoil; diesel delivery calls for more distribution centers and electric delivery for fewer. In a reference network that finances its stock in transit, slower line haul delivers nearly all of the cut in emissions, and the network grows mainly for deeper cuts or at high carbon prices. Several documented moves

in practice are consistent with these directions; in the sources we examine, firms report them in terms of cost, delivery service, and transport miles. Our model adds the carbon of the inventory a slower chain holds, the financing of the goods it carries, and the link between speed and density.

Limitations. The model describes an average, symmetric distribution environment, so its classification serves as a diagnostic for a heterogeneous network. Heterogeneous markets or discrete facility sites replace $d(n) = \delta n^{-1/2}$ with the realized last-mile profile; this can move the density threshold but leaves the comparison between facility and last-mile carbon per dollar intact. The representative inbound distance isolates the proximity effect of adding distribution centers. If more centers also mean smaller inbound shipments, lower vehicle fill, or lost facility scale, those forces favor consolidation and can move a case near the threshold, as last-mile fill does (Section 5.3). We charge the financing of goods in transit to the firm that owns them; a firm that buys on delivered terms may still bear that cost through the supplier's price. We treat speed as a continuous index between real options on a lane, and we assume that all inbound volume can use each option; where only part of it can move by rail, slowing reaches a smaller share of emissions. Carload transit times also vary more than intermodal's (Kwon et al. 1995), and the model, which treats transit time as fixed, does not price that variability. A firm that splits volume between modes faces a mixture of lead times, and our safety-stock expression does not represent such a mixture. The model ties last-mile distance to the number of stocking distribution centers, whereas networks that deliver through cross-docks or hubs separate stocking density from delivery density. Service enters as a cycle-service requirement; if it is chosen, safety stock becomes a further design margin. Finally, we calibrate the carbon intensities to public industry data, so a firm should compute the classification with its own rates, and the magnitudes in Section 5 describe the network we calibrate.

Future research. Future research could extend the model to multi-product and multi-echelon settings, where a shared network serves products in different regimes and safety stock can be held at more than one echelon. Separating stocking from delivery density, and modeling shipment consolidation and routing explicitly, would quantify the centralizing forces represented here by fill sensitivity. As accounting under ISO 14083 makes firm-level logistics-emissions data available (International Organization for Standardization 2023), empirical work can locate real products relative to the speed and density thresholds. Endogenous service choice and explicit mode mixes would make the service target and a lane's mode split design decisions.

A firm that accounts for its distribution carbon must decide how fast to replenish and how many distribution centers to operate; each decision depends on which activity provides availability or proximity at the lower carbon intensity per dollar.

Data availability

The data and code supporting the numerical illustrations are available from the author on reasonable request.

References

Amazon.com, Inc. (2023). Second quarter 2023 earnings conference call. August 3, 2023.

Amazon.com, Inc. (2026). 5 things to know from Amazon's 2025 sustainability report.

American Transportation Research Institute (2025). An analysis of the operational costs of trucking: 2025 update. Technical report, American Transportation Research Institute.

Ansari, S., Başdere, M., Li, X., Ouyang, Y., and Smilowitz, K. (2018). Advancements in continuous approximation models for logistics and transportation systems: 1996–2016. *Transportation Research Part B: Methodological*, 107:229–252.

Arıkan, E. and Jammernegg, W. (2014). The single-period inventory model under dual sourcing and product carbon footprint constraint. *International Journal of Production Economics*, 157:15–23.

Bektaş, T. and Laporte, G. (2011). The pollution-routing problem. *Transportation Research Part B: Methodological*, 45(8):1232–1250.

Benjaafar, S., Li, Y., and Daskin, M. (2013). Carbon footprint and the management of supply chains: Insights from simple models. *IEEE Transactions on Automation Science and Engineering*, 10(1):99–116.

Boone, T. and Ganeshan, R. (2025). Freight mode choice with emission caps: Revisiting classical inventory and transportation decisions. *Sustainability*, 17(9):4135.

Bouchery, Y., Ghaffari, A., Jemai, Z., and Dallery, Y. (2012). Including sustainability criteria into inventory models. *European Journal of Operational Research*, 222(2):229–240.

Cachon, G. P. (2014). Retail store density and the cost of greenhouse gas emissions. *Management Science*, 60(8):1907–1925.

Cariou, P. (2011). Is slow steaming a sustainable means of reducing CO2 emissions from container shipping? *Transportation Research Part D: Transport and Environment*, 16(3):260–264.

Chen, X., Benjaafar, S., and Elomri, A. (2013). The carbon-constrained EOQ. *Operations Research Letters*, 41(2):172–179.

Chen, X. and Wang, X. (2016). Effects of carbon emission reduction policies on transportation mode selections with stochastic demand. *Transportation Research Part E: Logistics and Transportation Review*, 90:196–205.

Comas Martí, J. M., Tancrez, J.-S., and Seifert, R. W. (2015). Carbon footprint and responsiveness trade-offs in supply chain network design. *International Journal of Production Economics*, 166:129–142.

Cushman & Wakefield (2025). U.S. industrial MarketBeat, Q2 2025.

Daganzo, C. F. (2005). *Logistics Systems Analysis*. Springer, 4th edition.

Daskin, M. S., Coullard, C. R., and Shen, Z.-J. M. (2002). An inventory-location model: Formulation, solution algorithm and computational results. *Annals of Operations Research*, 110:83–106.

Drent, M., Moradi, P., and Arts, J. (2023). Efficient emission reduction through dynamic supply mode selection. *European Journal of Operational Research*, 311(3):925–941.

Eppen, G. D. (1979). Effects of centralization on expected costs in a multi-location newsboy problem. *Management Science*, 25(5):498–501.

European Commission (2016). Commission delegated regulation (EU) 2016/2071 amending regulation (EU) 2015/757 as regards the methods for monitoring carbon dioxide emissions.

European Environment Agency (2026). EU new emissions trading system (ETS2) and social climate fund: Mechanisms, challenges and conditions for a fair decarbonisation in road transport. Technical report, European Environment Agency.

European Parliament and Council of the European Union (2023). Directive (EU) 2023/959 amending directive 2003/87/EC establishing a system for greenhouse gas emission allowance trading within the union. Official Journal of the European Union.

Federal Railroad Administration (2013). Higher speed freight truck market analysis. Technical Report DOT/FRA/ORD-13/32, U.S. Department of Transportation.

FreightWaves (2026). J.B. Hunt sees 'massive opportunities' amid intermodal shift.

Hacardiaux, T., Defryn, C., Tancrez, J.-S., and Verdonck, L. (2022). Balancing partner preferences for logistics costs and carbon footprint in a horizontal cooperation. *OR Spectrum*, 44(1):121–153.

Hacardiaux, T. and Tancrez, J.-S. (2020). Assessing the environmental benefits of horizontal cooperation using a location-inventory model. *Central European Journal of Operations Research*, 28(4):1363–1387.

Hammami, R., Nouira, I., and Frein, Y. (2015). Carbon emissions in a multi-echelon production-inventory model with lead time constraints. *International Journal of Production Economics*, 164:292–307.

Harris, I., Naim, M., Palmer, A., Potter, A., and Mumford, C. (2011). Assessing the impact of cost optimization based on infrastructure modelling on CO2 emissions. *International Journal of Production Economics*, 131(1):313–321.

Hoen, K. M. R., Tan, T., Fransoo, J. C., and van Houtum, G. J. (2014). Effect of carbon emission regulations on transport mode selection under stochastic demand. *Flexible Services and Manufacturing Journal*, 26(1–2):170–195.

Hu, H., Guo, S., and Shu, J. (2025). A location-inventory problem with facility temperature selection and quality deterioration in fresh product supply chain. *Transportation Research Part E: Logistics and Transportation Review*, 197:104044.

Hua, G., Cheng, T. C. E., and Wang, S. (2011). Managing carbon footprints in inventory management. *International Journal of Production Economics*, 132(2):178–185.

ICF International (2009). Comparative evaluation of rail and truck fuel efficiency on competitive corridors. Final report, Federal Railroad Administration, U.S. Department of Transportation.

International Organization for Standardization (2023). ISO 14083:2023 greenhouse gases—quantification and reporting of greenhouse gas emissions arising from transport chain operations.

Jalal, A., Adulyasak, Y., Jans, R., Morabito, R., and Toso, E. (2026). An integrated location-

inventory-transportation problem under demand uncertainty. *Transportation Research Part E: Logistics and Transportation Review*, 213:104990.

Jensen, J. L. W. V. (1906). Sur les fonctions convexes et les inégalités entre les valeurs moyennes. *Acta Mathematica*, 30:175–193.

Koca, E. (2021). Stochastic lot sizing problem with carbon emission constraints. *Gazi University Journal of Science*, 34(1):148–160.

Konur, D. and Schaefer, B. (2014). Integrated inventory control and transportation decisions under carbon emissions regulations: LTL vs. TL carriers. *Transportation Research Part E: Logistics and Transportation Review*, 68:14–38.

Kwon, O. K., Martland, C. D., Sussman, J. M., and Little, P. (1995). Origin-to-destination trip times and reliability of rail freight services in North American railroads. *Transportation Research Record*, 1489:1–8.

Li, S., Li, X., Zhang, D., and Zhou, L. (2017). Joint optimization of distribution network design and two-echelon inventory control with stochastic demand and CO2 emission tax charges. *PLoS ONE*, 12(1):e0168526.

Lineage, Inc. (2024). Prospectus (Form 424B4). U.S. Securities and Exchange Commission.

Maloni, M., Paul, J. A., and Gligor, D. M. (2013). Slow steaming impacts on ocean carriers and shippers. *Maritime Economics & Logistics*, 15(2):151–171.

McKinnon, A. C. and Piecyk, M. I. (2012). Setting targets for reducing carbon emissions from logistics: Current practice and guiding principles. *Carbon Management*, 3(6):629–639.

McKinsey & Company (2024). Decarbonizing logistics: Charting the path ahead.

Moradi, P., Arts, J., and Velázquez Martínez, J. C. (2025). Load asymptotics and dynamic speed optimization for the greenest path problem: A comprehensive analysis. *OR Spectrum*, 47(2):477–524.

Mulholland, E. (2022). Cost of electric commercial vans and pickup trucks in the United States through 2040. Working Paper 2022-01, International Council on Clean Transportation.

Ni, W. and Shu, J. (2015). Trade-off between service time and carbon emissions for safety stock placement in multi-echelon supply chains. *International Journal of Production Research*, 53(22):6701–6718.

Ouyang, Y. and Daganzo, C. F. (2006). Discretization and validation of the continuum approximation scheme for terminal system design. *Transportation Science*, 40(1):89–98.

Özsen, L., Coullard, C. R., and Daskin, M. S. (2008). Capacitated warehouse location model with risk pooling. *Naval Research Logistics*, 55(4):295–312.

Pharmaceutical Commerce (2023). LogiPharma USA 2023: Sea vs. air—identifying ways to make international cold shipping efficient and sustainable. October 10, 2023.

Psaraftis, H. N. and Kontovas, C. A. (2013). Speed models for energy-efficient maritime transportation: A taxonomy and survey. *Transportation Research Part C: Emerging Technologies*,

26:331–351.

Puttemans, I., Caris, A., and Braekers, K. (2025). The integration of location and inventory decisions in supply chain networks: A literature review and future prospects. *Transportation Research Part E: Logistics and Transportation Review*, 195:103970.

Qi, Y., Harrod, S., Psaraftis, H. N., and Lang, M. (2022). Transport service selection and routing with carbon emissions and inventory costs consideration in the context of the Belt and Road Initiative. *Transportation Research Part E: Logistics and Transportation Review*, 159:102630.

Schaefer, B. and Konur, D. (2015). Economic and environmental considerations in a continuous review inventory control system with integrated transportation decisions. *Transportation Research Part E: Logistics and Transportation Review*, 80:142–165.

Schmitt, A. J., Sun, S. A., Snyder, L. V., and Shen, Z.-J. M. (2015). Centralization versus decentralization: Risk pooling, risk diversification, and supply chain disruptions. *Omega*, 52:201–212.

Sebatjane, M. (2025). Sustainable inventory models under carbon emissions regulations: Taxonomy and literature review. *Computers & Operations Research*, 175:106865.

Shaw, K., Irfan, M., Shankar, R., and Yadav, S. S. (2016). Low carbon chance constrained supply chain network design problem: A benders decomposition based approach. *Computers & Industrial Engineering*, 98:483–497.

Shen, Z.-J. M., Coullard, C., and Daskin, M. S. (2003). A joint location-inventory model. *Transportation Science*, 37(1):40–55.

Ship & Bunker (2026). Analysis: Average VLSFO bunker prices at key ports fall 20% in 2025.

Silbermayr, L., Jammernegg, W., and Kischka, P. (2017). Inventory pooling with environmental constraints using copulas. *European Journal of Operational Research*, 263(2):479–492.

Snyder, L. V., Daskin, M. S., and Teo, C.-P. (2007). The stochastic location model with risk pooling. *European Journal of Operational Research*, 179(3):1221–1238.

State of California (2023). Climate corporate data accountability act (Senate Bill 253).

Straubert, C. (2024). A continuous approximation location-inventory model with exact inventory costs and nonlinear delivery lead time penalties. *International Journal of Production Economics*, 268:109092.

The Home Depot (2020). Supply chain unveils first flatbed distribution center (FDC). Company newsroom, January 28, 2020.

Tsao, Y.-C., Mangotra, D., Lu, J.-C., and Dong, M. (2012). A continuous approximation approach for the integrated facility-inventory allocation problem. *European Journal of Operational Research*, 222(2):216–228.

Turken, N., Carrillo, J., and Verter, V. (2017). Facility location and capacity acquisition under carbon tax and emissions limits: To centralize or to decentralize? *International Journal of Production Economics*, 187:126–141.

U.S. Energy Information Administration (2025). Electric power annual.

U.S. Energy Information Administration (2026). Weekly petroleum status report.

U.S. Environmental Protection Agency (2022). Emission factors for greenhouse gas inventories. Technical report, U.S. Environmental Protection Agency, Center for Corporate Climate Leadership.

U.S. Environmental Protection Agency (2024a). Supply chain greenhouse gas emission factors v1.3 by NAICS-6. Technical report, U.S. Environmental Protection Agency.

U.S. Environmental Protection Agency (2024b). U.S. energy use intensity by property type: U.S. national median reference values. Technical report, ENERGY STAR Portfolio Manager.

Wang, M., Wu, J., Kafa, N., and Klibi, W. (2020). Carbon emission-compliance green location-inventory problem with demand and carbon price uncertainties. *Transportation Research Part E: Logistics and Transportation Review*, 142:102038.

Wang, Y., Shu, J., Su, J., and Zhen, L. (2025). Inventory placement with carbon cap-and-trade in guaranteed service supply chains. *Transportation Research Part E: Logistics and Transportation Review*, 193:103813.

World Bank (2021). Decarbonizing the freight and logistics sector: Discussion paper. Transport decarbonization investment series, World Bank Group.

Zhang, J., Yang, L., and Campbell, J. F. (2026). Economic and environmental impacts of integrated drone delivery systems. *Transportation Research Part E: Logistics and Transportation Review*, 209:104759.

A Proofs

Proof of Lemma 4.1. We give the argument for J; for C replace each blended rate by its carbon rate and set $h_p = 0$. Let $x = \ln n$ and $y = \ln v$, which range over a box, and let $k = z\sigma\sqrt{ND}$. Then

$$J = F_\tau e^x + H_\tau k\, e^{g(x,y)} + h_p N\mu D\, e^{-y} + T_\tau(e^y)\, N\mu D + \ell_\tau N\mu\delta\, e^{-x/2},$$

with $g(x,y) = \frac{1}{2}\ln\big(N\rho + (1-\rho)e^x\big) - \frac{1}{2}y$. (i) The facility and last-mile terms are strictly convex in x. The function g is convex, because $\ln(c_1 + c_2 e^x)$ is convex for $c_1 \geq 0$ and $c_2 > 0$, so the holding term is log-convex and therefore convex. The transport term is convex in y, because T_τ is convex and nondecreasing and e^y is convex. For strict convexity, let $S = H_\tau k\, e^g$; its Hessian is $S\big(\nabla^2 g + \nabla g\, \nabla g^\mathsf{T}\big)$, with $\nabla g = (g_x, -\frac{1}{2})$ and $\nabla^2 g = \mathrm{diag}(g_{xx}, 0)$, $g_{xx} \geq 0$. The Hessian of the facility, holding, and last-mile terms therefore has (x,x) entry $F_\tau e^x + \frac{1}{4}\ell_\tau N\mu\delta e^{-x/2} + S(g_{xx} + g_x^2) > 0$, (y,y) entry $S/4$, and off-diagonal entry $-Sg_x/2$, so its determinant is $\big(F_\tau e^x + \frac{1}{4}\ell_\tau N\mu\delta e^{-x/2}\big)S/4 + S^2 g_{xx}/4 > 0$. These terms are strictly convex, and adding the convex transport and financing terms preserves strict convexity; the financing term $h_p N\mu D\, e^{-y}$ is convex in y and does not depend on x. (ii) A strictly convex continuous function has a unique minimizer on a compact convex set. Since J is continuous in (x, y, τ), the minimizer is continuous in τ by the maximum theorem. (iii) For fixed x,

J is strictly convex in y, so the speed response is unique and, by the maximum theorem, continuous. The transport, financing, facility, and last-mile terms are additively separable in x and y, and the cross-partial of the holding term is $-Sg_x/2 \le 0$, so J has decreasing differences in (x, y), and the minimizing y is nondecreasing in x and the minimizing x nondecreasing in y. At a fixed speed, J is strictly convex in x; the reduced objective is the partial minimum of a strictly convex function over a convex set and is therefore strictly convex in x. A strictly convex function of $\ln n$ is strictly quasiconvex in n. By Danskin's theorem, which applies because the speed response is unique, J^{r} is differentiable with $J^{\mathrm{r}}_n(n) = J_n\big(n, v^*(n)\big)$, and by strict quasiconvexity J^{r}_n changes sign at most once, from negative to positive. For strict monotonicity with $\rho < 1$: where v^* is interior, (12) holds; its right side strictly increases in n, through SS, and strictly decreases in v, while its left side does not depend on n and is nondecreasing in v, so v^* strictly increases in n, and projection onto the speed bounds preserves monotonicity. Since $G(n, v) \propto v^{-1/2}$, $J_{nv} = -H_\tau G(n, v)/(2v) < 0$, and $J_n(\cdot, v)$ changes sign at most once, from negative to positive, so its sign change occurs at a higher n when v rises. □

Proof of Proposition 4.2. (i) $N\mu D\, t(v)$, $N\mu D\, \varepsilon(v)$, and $h_p N\mu D\, v^{-1}$ are convex and $\kappa(n)v^{-1/2}$ has $\partial^2_v = \frac{3}{4}\kappa(n)v^{-5/2} > 0$, so both objectives are strictly convex in v. We use $M_s = \frac{1}{2}h\,\kappa(n)\, v_{\mathrm{K}}^{-3/2} = h\mathrm{SS}/(2v_{\mathrm{K}})$ and $M_p = h_p N\mu D\, v_{\mathrm{K}}^{-2} = h_p P/v_{\mathrm{K}}$ to denote the marginal savings of safety stock and of stock in transit, so that $a(n, v_{\mathrm{K}}) = M_s/(M_s + M_p)$. The cost optimum satisfies $\hat{K}'(v_{\mathrm{K}}) = 0$, i.e. $N\mu D\, t'(v_{\mathrm{K}}) = M_s + M_p$, so that $\frac{1}{2}\kappa(n)v_{\mathrm{K}}^{-3/2} = a(n, v_{\mathrm{K}})\, N\mu D\, t'(v_{\mathrm{K}})/h$. Evaluating the carbon derivative there,

$$\hat{C}'(v_{\mathrm{K}}) = -\eta\, \tfrac{1}{2}\kappa(n)v_{\mathrm{K}}^{-3/2} + N\mu D\, \varepsilon'(v_{\mathrm{K}}) = N\mu D\big(\varepsilon'(v_{\mathrm{K}}) - a(n, v_{\mathrm{K}})\, (\eta/h)\, t'(v_{\mathrm{K}})\big),$$

which is negative — so $v_C > v_{\mathrm{K}}$ by strict convexity — if and only if (13). With $t' = 2t_1 v$ and $\varepsilon' = 2\varepsilon_0 v$ the marginal ratio is ε_0/t_1, and with $h_p = 0$ the priced first-order condition (12) gives $v^{5/2} = (h+\tau\eta)z\sigma\sqrt{ND}\,\sqrt{\Psi}/[4(t_1+\tau\varepsilon_0)N\mu D]$, which is (14). The projection and the boundary qualification are immediate. (ii) Let $s = z\sigma$, so that $\mathrm{SS} \propto s\, v^{-1/2}$. Under the quadratic instantiation, condition (12) takes the form $M_s + M_p = 2(t_1 + \tau\varepsilon_0)N\mu D\, v$ with $M_s = H_\tau \mathrm{SS}/(2v) \propto s\, v^{-3/2}$ and $M_p = h_p N\mu D\, v^{-2}$; let $a_\tau = M_s/(M_s + M_p)$. Implicit differentiation gives $\mathrm{d}\ln v^*/\mathrm{d}\ln s = M_s/(\frac{5}{2}M_s + 3M_p) = a_\tau/(3 - a_\tau/2) \in (0, \frac{2}{5}]$, so speed rises with service and demand volatility; without financing $v^* \propto s^{2/5}$, as in (14). At $\tau = 0$, $M_s/M_p \propto s\, v_{\mathrm{K}}^{1/2}$ rises in s, directly and through v_{K}, so $a(n, v_{\mathrm{K}})$ rises in s; the right side of (13) is the constant ε_0/t_1, so the comparison can change only toward faster replenishment, and without financing it does not change. □

Proof of Proposition 4.3. (i) By Lemma 4.1(iii), $J(\cdot, v)$ and $C(\cdot, v)$ are strictly quasiconvex on $[1, N]$, which gives uniqueness and the interior/boundary characterization. For the threshold, take $\tau = 0$ for cost and $\tau \to \infty$ for carbon. The cost density satisfies $\partial_n K = F + h\, G(n_K, v) - t_L N\mu\delta/(2n_K^{3/2}) = 0$, so the last-mile marginal $t_L N\mu\delta/(2n_K^{3/2}) = F + h\, G(n_K, v)$. Evaluating the

carbon derivative there,

$$\partial_n C\big|_{n_K} = \varphi + \eta\, G(n_K, v) - \frac{\varepsilon_L N \mu \delta}{2 n_K^{3/2}} = \varphi + \eta\, G(n_K, v) - \frac{\varepsilon_L}{t_L}\big[F + h\, G(n_K, v)\big],$$

which is negative — so $n_C > n_K$ by strict quasiconvexity of $C(\cdot, v)$ — if and only if (16). With $R(G) = (\varphi + \eta G)/(F + hG)$, we have $R(0) = \varphi/F$, $R(G) \to \eta/h$ as $G \to \infty$, and $R'(G) = (\eta F - h\varphi)/(F + hG)^2$, so R is monotone with $\operatorname{sign} R' = \operatorname{sign}(\eta/h - \varphi/F)$ and lies between φ/F and η/h. In the no-pooling limit $\rho \to 1$, $G \to 0$ and $\Psi \to N$, so (11) gives $n^* = (\ell_\tau N \mu \delta / 2F_\tau)^{2/3}$ and (16) becomes (17). (ii) Along the speed response of Proposition 4.2(ii), $J^{\mathrm{r}}_n(n; s) = J_n\big(n, v^*(n; s); s\big)$ because $J_v = 0$. The holding marginal term in (11) is proportional to $s(v^*)^{-1/2}(1-\rho)/\sqrt{\Psi(n)}$, whose elasticity in s is $1 - a_\tau/(6 - a_\tau) \geq \frac{4}{5}$, so it rises in s even after the induced speed increase, and $J^{\mathrm{r}}_{ns} > 0$. By Lemma 4.1(iii), $J^{\mathrm{r}}_n(\cdot\,; s)$ changes sign at most once, from negative to positive, and it is increasing in s; its sign change therefore occurs at a lower n when s rises, so the optimal density is nonincreasing in s, and strictly decreasing where it is interior. □

Proof of Proposition 4.4. (i) By Lemma 4.1 applied to the carbon objective, $n_C > n_K$ if and only if the reduced carbon objective slopes downward at n_K, $C^{\mathrm{r}}_n(n_K) < 0$, and $n_C = n_K$ if it is zero. By the envelope theorem (Danskin's theorem, since $\hat{v}_C$ is unique whether or not it is at a bound), $C^{\mathrm{r}}_n(n_K) = \varphi + \eta\, G(n_K, \hat{v}_C) - \frac{1}{2}\varepsilon_L N \mu \delta\, n_K^{-3/2}$. The financing term does not depend on n, so the cost condition (11) at $\tau = 0$, which holds because n_K is interior, gives $\frac{1}{2} t_L N \mu \delta\, n_K^{-3/2} = F + hG(n_K, v_K)$, and the last term equals $r_L\big[F + hG(n_K, v_K)\big]$. Since $G \propto v^{-1/2}$, $\eta\, G(n_K, \hat{v}_C) = h\, G(n_K, v_K)\, r_h \sqrt{v_K/\hat{v}_C}$. Dividing $C^{\mathrm{r}}_n(n_K)$ by $F + hG(n_K, v_K) > 0$ and substituting $\varphi = F\, r_F$ gives (20). Under the quadratic instantiation with v_K and $\hat{v}_C$ interior, the carbon condition (12) at n_K gives $\eta \mathrm{SS}(n_K, \hat{v}_C) = 4\varepsilon_0 N \mu D\, \hat{v}_C^2$, and the cost condition gives $h\mathrm{SS}(n_K, v_K) = 4 a_K t_1 N \mu D\, v_K^2$; with SS $\propto v^{-1/2}$ these imply $(\hat{v}_C/v_K)^{5/2} = a_K R_{\mathrm{speed}}$, so $r_h\sqrt{v_K/\hat{v}_C} = r_h^{4/5} r_t^{1/5} a_K^{-1/5}$, which is (21). The second expression for ω in (18) follows from $F + hG(n_K, v_K) = \frac{1}{2} t_L N \mu \delta\, n_K^{-3/2}$ and $F = \frac{1}{2} t_L N \mu \delta\, (n_K^0)^{-3/2}$. (ii) The carbon objective has no financing term, so by (12) its speed response satisfies $v_C(n)^{5/2} \propto \sqrt{\Psi(n)}$, and $\ln[v_C/\hat{v}_C] = \frac{1}{5}\ln[\Psi(n_C)/\Psi(n_K)]$. Adding $\ln[\hat{v}_C/v_K] = \frac{2}{5}\ln(a_K R_{\mathrm{speed}})$ from part (i) gives (22). If (21) holds and $a_K R_{\mathrm{speed}} \geq 1$, then $n_C > n_K$ and $\Psi(n_C) > \Psi(n_K)$, so in (22) the first term is nonnegative and the second positive; the reversed case is symmetric. Since Ψ is increasing with $\Psi(1) \leq \Psi(n_C) \leq \Psi(N) = N$, the threshold $\sqrt{\Psi(n_K)/\Psi(n_C)}$ lies in $[\underline{R}, \overline{R}]$. (iii) With $\rho = 1$, $\Psi \equiv N$ and $G \equiv 0$, so (11) involves only F_τ and ℓ_τ and gives $n^*(\tau) = (\ell_\tau N \mu \delta / 2F_\tau)^{2/3}$; the ratio of its carbon and cost values is $(r_L/r_F)^{2/3}$, and the financing term does not depend on n. The speed response no longer depends on n, and by the argument of part (i) the ratio of its carbon and cost values is $(a_K R_{\mathrm{speed}})^{2/5}$. For (23), carbon separates into a proximity component $C_{\mathrm{prox}}(n) = \varphi n + \varepsilon_L N \mu \delta n^{-1/2}$ and an availability component $C_{\mathrm{avail}}(v) = \eta z \sigma N \sqrt{D/v} + \varepsilon_0 N \mu D\, v^2$, with $\theta = C_{\mathrm{prox}}(n_C)/C_{\min}$. A function $c_1 n + c_2 n^{-1/2}$ attains its minimum $3c_1 n^\circ$ at $n^\circ = (c_2/2c_1)^{2/3}$ and equals $c_1 n^\circ(x + 2x^{-1/2})$ at $n = x n^\circ$; with $n_K = x n_C$ and $x = R_{\mathrm{density}}^{-2/3}$, $C_{\mathrm{prox}}(n_K)/C_{\mathrm{prox}}(n_C) = (x + 2x^{-1/2})/3$. Likewise $c_1 v^{-1/2} + c_2 v^2$ attains its minimum $\frac{5}{4} c_1 (v^\circ)^{-1/2}$ and equals $c_1 (v^\circ)^{-1/2}(y^{-1/2} + y^2/4)$ at $v = y v^\circ$, so with $v_K = y v_C$ and $y = (a_K R_{\mathrm{speed}})^{-2/5}$,

$C_{\text{avail}}(v_K)/C_{\text{avail}}(v_C) = (4y^{-1/2} + y^2)/5$. Weighting by θ and $1-\theta$ gives $C(n_K, v_K)/C_{\min}$ exactly, and expanding $x = e^{u_1}$ and $y = e^{u_2}$ to second order gives (23). □

Proof of Proposition 4.5. (i) Condition (12) is $\Upsilon(v,\tau) = 2b_\tau N\mu D\, v - M_s - M_p = 0$, with $b_\tau = t_1 + \tau\varepsilon_0$, $M_s = H_\tau \mathrm{SS}/(2v)$, and $M_p = h_p P/v$. Then $v\,\Upsilon_v = \frac{5}{2}M_s + 3M_p$ and, since $\tau\eta/H_\tau = s_h$ and $\tau\varepsilon_0/b_\tau = s_t$, $\tau\,\Upsilon_\tau = s_t(M_s + M_p) - s_h M_s$. Hence $\mathrm{d}\ln v^*/\mathrm{d}\ln\tau = -\tau\Upsilon_\tau/(v\Upsilon_v) = (a_\tau s_h - s_t)/(3 - a_\tau/2)$ with $a_\tau = M_s/(M_s + M_p)$. Substituting the definitions, $a_\tau s_h - s_t = \tau\,(h\mathrm{SS} + 2h_pP)\,t_1\,[a(n,v^*)\,r_h - r_t]/[(H_\tau \mathrm{SS} + 2h_pP)\,b_\tau]$, which has the sign of $a(n,v^*)r_h - r_t$. Because $a(n,v)$ rises with v, a positive sign raises v^* and keeps the sign positive, and a negative sign lowers v^* and keeps it negative, so the sign does not change along the path and v^* is monotone in τ. With $h_p = 0$, $a_\tau = 1$ and the elasticity is $\frac{2}{5}(s_h - s_t)$. (ii) The same computation applies to $n^* = (\ell_\tau N\mu\delta/2F_\tau)^{2/3}$. (iii) With $h_p = 0$, by (14), along the speed response $H_\tau G\big(n, v^*(n)\big) = c\,(1-\rho)\,H_\tau^{4/5} b_\tau^{1/5}\Psi(n)^{-3/5}$, with $b_\tau = t_1 + \tau\varepsilon_0$ and $c > 0$ depending on neither n nor τ. By the envelope theorem, $J_n^{\mathrm{r}} = F_\tau + H_\tau G\big(n, v^*(n)\big) - \frac{1}{2}\ell_\tau N\mu\delta n^{-3/2}$; multiplying by $n^{3/2}/\ell_\tau > 0$ preserves its sign,

$$\frac{n^{3/2} J_n^{\mathrm{r}}}{\ell_\tau} = \frac{F_\tau}{\ell_\tau}\, n^{3/2} + c\,(1-\rho)\,\frac{H_\tau^{4/5} b_\tau^{1/5}}{\ell_\tau}\,\frac{n^{3/2}}{\Psi(n)^{3/5}} - \tfrac{1}{2} N\mu\delta.$$

This expression changes sign at most once, from negative to positive (Lemma 4.1(iii)). If both coefficients are nonincreasing in τ, the expression is pointwise nonincreasing in τ, so its sign change moves weakly to higher n and $n^*(\tau)$ is nondecreasing. The logarithmic derivatives of the coefficients are $1/(r_F^{-1} + \tau) - 1/(r_L^{-1} + \tau)$ and $\frac{4}{5}/(r_h^{-1} + \tau) + \frac{1}{5}/(r_t^{-1} + \tau) - 1/(r_L^{-1} + \tau)$. Since $1/(r^{-1} + \tau)$ is increasing in r, both are nonpositive at every τ when r_L exceeds r_F, r_h, and r_t, and nonnegative when it falls short of each. □

Proof of Corollary 4.6. (i) In (24) the throughput coefficient $\gamma(n,v) = \varepsilon(v)D + \varepsilon_L d(n) > 0$, so $\{\tilde{C} \le B\} = \{Q \le (B - \varphi n - \eta\mathrm{SS})/\gamma\}$, and $\Pr\{\tilde{C} \le B\} \ge 1 - \beta_c$ holds iff $(B - \varphi n - \eta\mathrm{SS})/\gamma \ge q_{1-\beta_c}$, with $q_{1-\beta_c} = \inf\{q : \Pr(Q \le q) \ge 1 - \beta_c\}$, i.e. $\varphi n + \eta\mathrm{SS} + \gamma\, q_{1-\beta_c} \le B$. Because $\gamma\, q_{1-\beta_c} = (q_{1-\beta_c}/N\mu)\,[\varepsilon(v)N\mu D + \varepsilon_L N\mu d(n)]$, this condition is (25). Under normality $q_{1-\beta_c} = N\mu + z_{\beta_c}\sigma_Q$, so $m = 1 + z_{\beta_c}\,\mathrm{CV}_Q$. From $\sigma_Q^2 = \sigma^2 N[1 + (N-1)\rho]$, $\mathrm{CV}_Q = (\sigma/\mu)\sqrt{[1 + (N-1)\rho]/N}$ is increasing in σ/μ and, for $N > 1$, in ρ, and z_{β_c} is increasing in $1 - \beta_c$, so, for $\beta_c < \frac{1}{2}$, m is increasing in all three. (ii) The chance constraint multiplies the inbound and last-mile carbon rates by m; substituting $m\varepsilon_0$ and $m\varepsilon_L$ for ε_0 and ε_L in (13) and (16) gives the two conditions, the latter scaling only the left side. The same substitution in Proposition 4.4, with n_K, ω, and a_K unchanged because the cost rates are unchanged, gives the joint conditions. For the joint speed, divide the carbon density condition by m: its facility and lost-pooling terms then scale as m^{-1} and $m^{-4/5}$ and, multiplied by $n^{3/2}$, rise in n with elasticities $\frac{3}{2}$ and at least $\frac{9}{10}$, so that, at an interior n_C, $\mathrm{d}\ln n_C/\mathrm{d}\ln m \le \frac{8}{9}$. Since Ψ has elasticity at most one in n, (22) with $a_K R_{\text{speed}}/m$ in place of $a_K R_{\text{speed}}$ gives $\mathrm{d}\ln v_C/\mathrm{d}\ln m \le -\frac{2}{5} + \frac{1}{5}\cdot\frac{8}{9} < 0$. □